

# *Lecture 10: Training CNNs*

*Thursday March 2, 2017*

# Announcements!

- Final Project proposals due this **Today**
- I will be out of town next week. Rishit will lead class discussions.
- Next paper: **March 7** *You Only Look Once: Unified, Real-Time Object Detection*. If this paper seems too deep or confusing, look at *Fast R-CNN*, *Faster R-CNN*

# Opportunity: Google Brain Residency

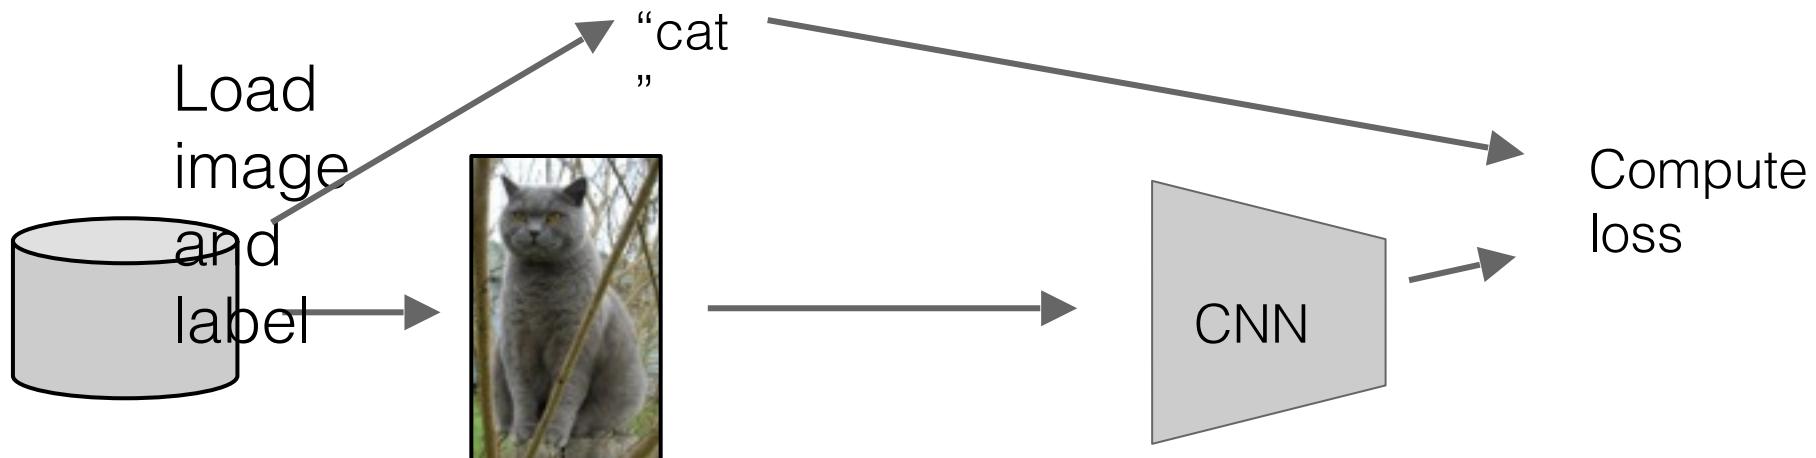
## What Is The Brain Residency Program?

The Google Brain Residency Program is a one-year intensive residency program focused on Deep Learning. Residents will have the opportunity to conduct cutting-edge research and work alongside some of the most distinguished deep learning scientists within the Google Brain team. To learn more about the team and what we do, visit [g.co/brain](http://g.co/brain)

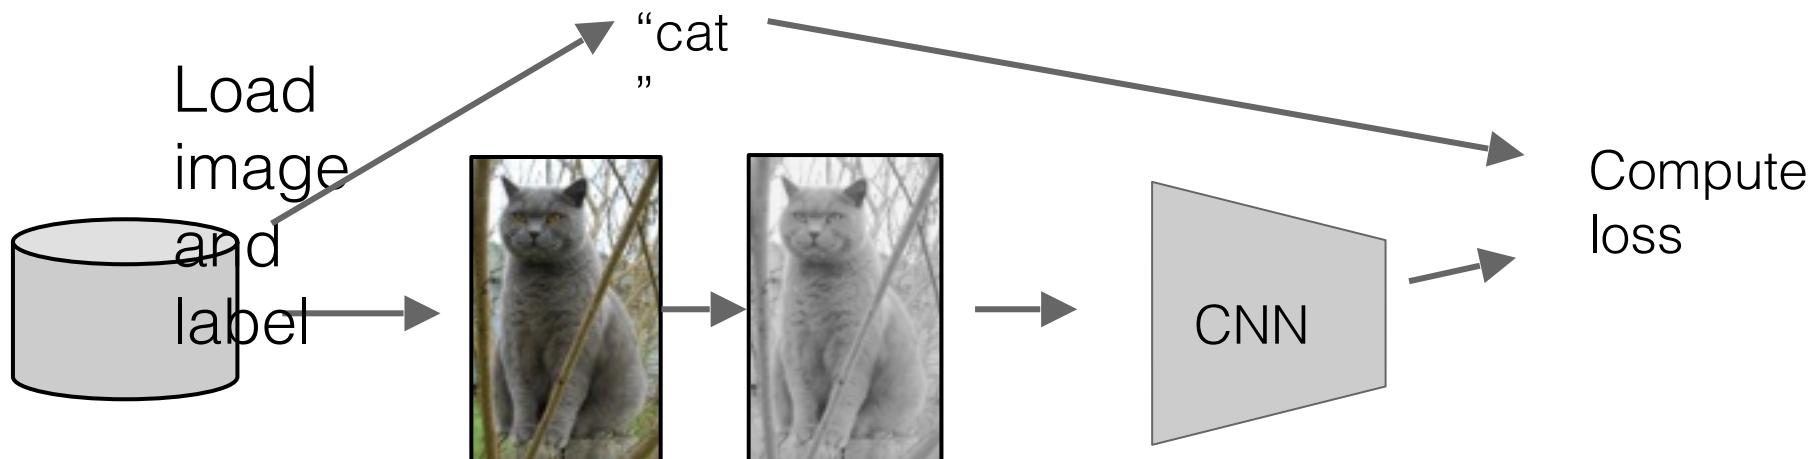
- Email contact for questions: [brain-residency@google.com](mailto:brain-residency@google.com)
- For more information on the Residency Program, check out our website at [g.co/brainresidency](http://g.co/brainresidency)
- More recently, we published a blog post on the Google Research Blog where we discuss updates on current Residents' progress and our program focus for 2017.

# Data Augmentation

# Data Augmentation

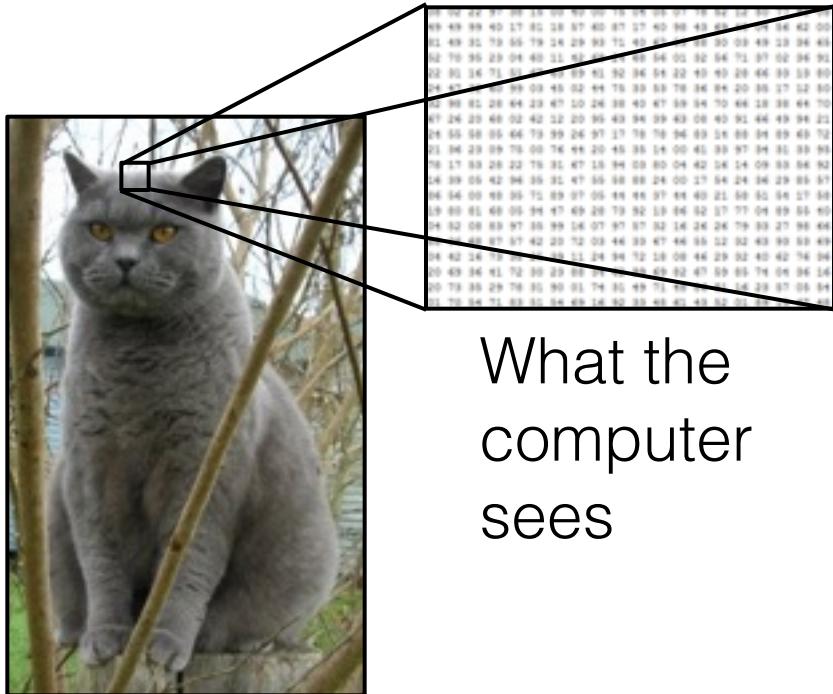


# Data Augmentation



# Data Augmentation

- Change the pixels without changing the label
- Train on transformed data
- VERY widely used



# Data Augmentation

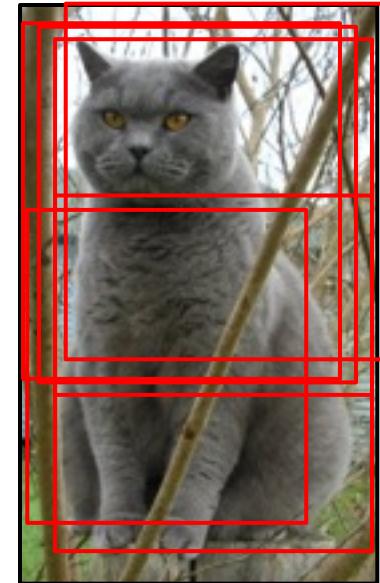
## 1. Horizontal flips



# Data Augmentation

## 2. Random crops/scales

**Training:** sample random crops / scales



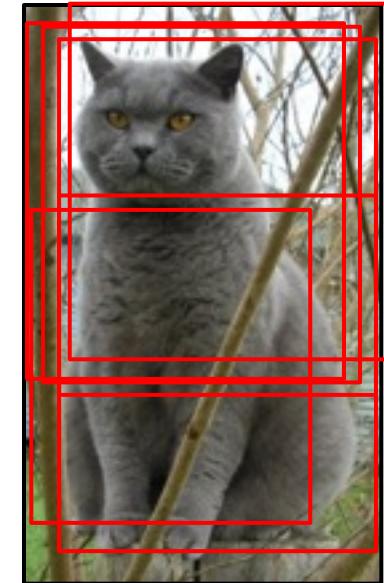
# Data Augmentation

## 2. Random crops/scales

**Training:** sample random crops / scales

ResNet:

1. Pick random  $L$  in range  $[256, 480]$
2. Resize training image, short side =  $L$
3. Sample random  $224 \times 224$  patch



# Data Augmentation

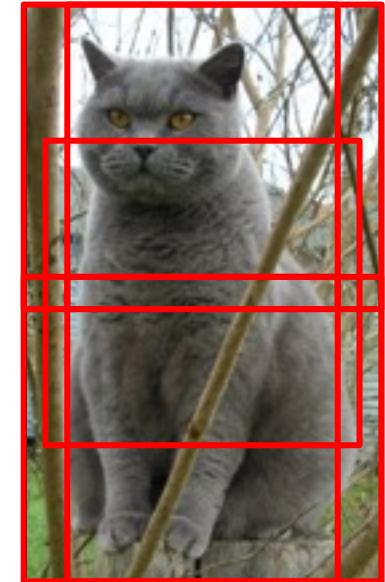
## 2. Random crops/scales

**Training:** sample random crops / scales

ResNet:

1. Pick random  $L$  in range [256, 480]
2. Resize training image, short side =  $L$
3. Sample random  $224 \times 224$  patch

**Testing:** average a fixed set of crops



# Data Augmentation

## 2. Random crops/scales

**Training:** sample random crops / scales

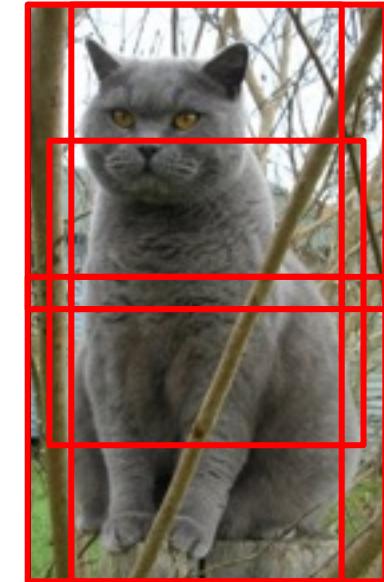
ResNet:

1. Pick random  $L$  in range [256, 480]
2. Resize training image, short side =  $L$
3. Sample random  $224 \times 224$  patch

**Testing:** average a fixed set of crops

ResNet:

1. Resize image at 5 scales: {224, 256, 384, 480, 640}
2. For each size, use 10  $224 \times 224$  crops: 4 corners + center, + flips



# Data Augmentation

## 3. Color jitter

**Simple:**

Randomly jitter contrast



# Data Augmentation

## 3. Color jitter

**Simple:**

Randomly jitter contrast



**Complex:**

1. Apply PCA to all [R, G, B] pixels in training set
2. Sample a “color offset” along principal component directions
  1. Add offset to all pixels of a training image  
(As seen in *[Krizhevsky et al. 2012]*, ResNet, etc)

# Data Augmentation

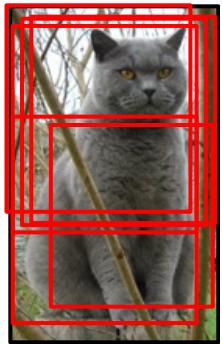
## 3. Color jitter

Random mix/combinations of :

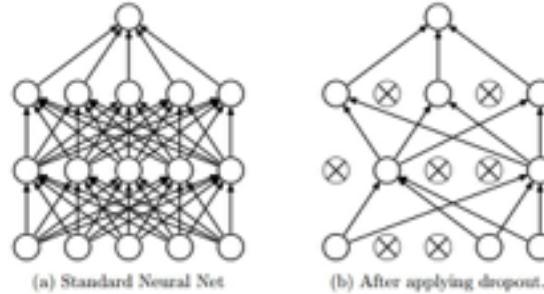
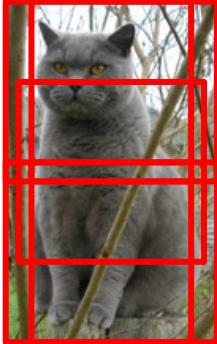
- translation
- rotation
- stretching
- shearing,
- lens distortions, ... (go crazy)

# A general theme:

1. **Training:** Add random noise
2. **Testing:** Marginalize over the noise



Data Augmentation



Dropout

Batch normalization,  
Model ensembles

# Data Augmentation: Takeaway

- Simple to implement, use it
- Especially useful for small datasets
- Fits into framework of noise / marginalization

# Transfer Learning

“You need a lot of data if you want  
to train/use CNNs”

# Transfer Learning

“You need a lot of data if you want  
to train/use CNNs”

**Not True**

# Transfer Learning with CNNs



1. Train on  
Imagenet

# Transfer Learning with CNNs



1. Train on  
Imagenet

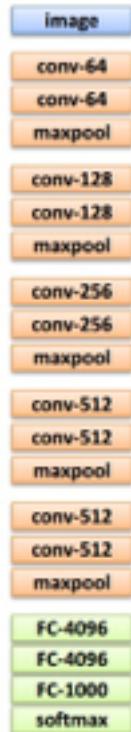


2. Small  
dataset:  
feature  
extractor

Freeze  
these

Train  
this

# Transfer Learning with CNNs



1. Train on  
Imagenet



2. Small  
dataset:  
**feature  
extractor**

Freeze  
these

Train  
this



3. Medium  
dataset:  
**finetuning**

Freeze  
these

more data = retrain  
more of the network (or  
all of it)

Train  
this

# Transfer Learning with CNNs



1. Train on  
Imagenet



2. Small  
dataset:  
**feature  
extractor**

Freeze  
these

Train  
this



3. Medium  
dataset:  
**finetuning**

Freeze  
these

more data = retrain  
more of the network (or  
all of it)

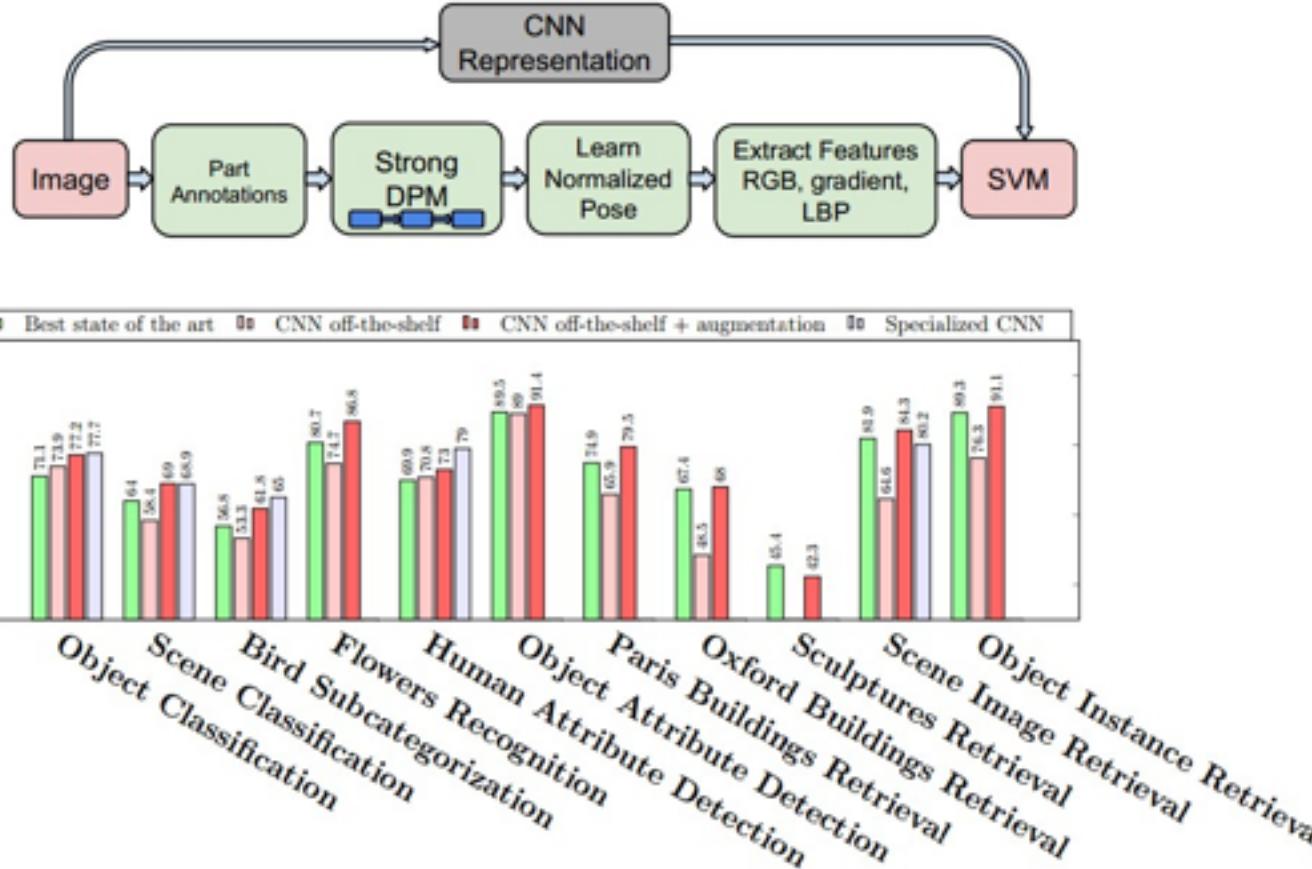
Train  
this

tip: use only ~1/10th of the original learning rate in finetuning top layer, and ~1/100th on intermediate layers

# CNN Features off-the-shelf: an Astounding Baseline for Recognition [Razavian et al, 2014]

DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition  
[Donahue\*, Jia\*, et al., 2013]

	DeCAF <sub>6</sub>	DeCAF <sub>7</sub>
LogReg	<b>40.94 ± 0.3</b>	40.84 ± 0.3
SVM	39.36 ± 0.3	40.66 ± 0.3
Xiao et al. (2010)	38.0	



\* Original slides borrowed from Andrej Karpathy and Li Fei-Fei, Stanford cs231n



	<b>very similar dataset</b>	<b>very different dataset</b>
<b>very little data</b>	?	?
<b>quite a lot of data</b>	?	?

image		
conv-64		
conv-64		
maxpool		
conv-128	more generic	
conv-128		
maxpool		
conv-256		
conv-256		
maxpool		
conv-512	more specific	
conv-512		
maxpool		
conv-512		
conv-512		
maxpool		
FC-4096		
FC-4096		
FC-1000		
softmax		

Diagram illustrating the layers of a neural network and their relationship to dataset similarity and training strategy:

- Image Input:** Represented by a blue box labeled "image".
- Convolutional Layers:** Represented by orange boxes. The first four layers (conv-64, conv-64, maxpool, conv-128) are grouped under the label "more generic". The next three layers (conv-256, conv-256, maxpool) are grouped under the label "more specific".
- FC Layers:** Represented by green boxes. These layers are at the bottom of the network.
- Dataset Similarity:** The columns represent different levels of dataset similarity:
  - very similar dataset:** Corresponds to the first two rows of layers.
  - very different dataset:** Corresponds to the last two rows of layers.
- Training Strategy:** The rightmost column provides recommendations based on dataset similarity:
  - very little data:** Use Linear Classifier on top layer.
  - quite a lot of data:** Finetune a few layers.
  - For "very different dataset", the strategy is marked with a question mark (?)

image		
conv-64		
conv-64		
maxpool		
conv-128	more generic	
conv-128		
maxpool		
conv-256	more specific	
conv-256		
maxpool		
conv-512		
conv-512		
maxpool		
conv-512		
conv-512		
maxpool		
FC-4096		
FC-4096		
FC-1000		
softmax		

Diagram illustrating the layers of a neural network and their relationship to dataset similarity:

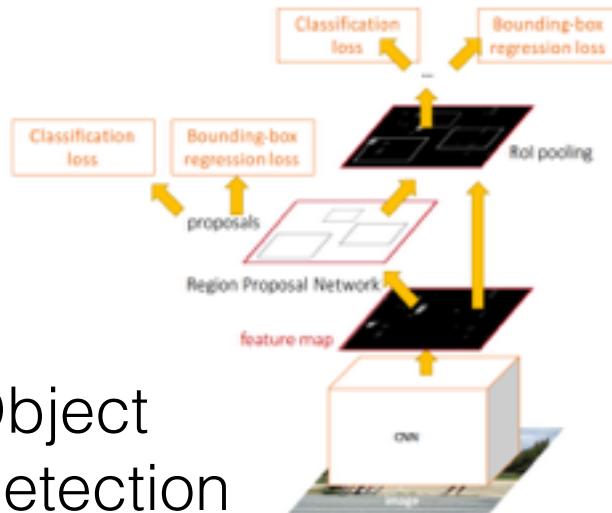
- Image Input:** Represented by a blue box labeled "image".
- Convolutional Layers:** Represented by orange boxes. They become "more generic" from top to bottom (conv-64, conv-64, maxpool) and then "more specific" (conv-128, conv-128, maxpool, conv-256, conv-256, maxpool, conv-512, conv-512, maxpool, conv-512, conv-512, maxpool).
- FC Layers:** Represented by green boxes (FC-4096, FC-4096, FC-1000, softmax).
- Dataset Similarity:** Indicated by arrows pointing from the layers to the columns of a table.
- Table Columns:**
  - very similar dataset:** Corresponds to the first two rows of the table.
  - very different dataset:** Corresponds to the last three rows of the table.

The table provides recommendations for training based on the amount of data available:

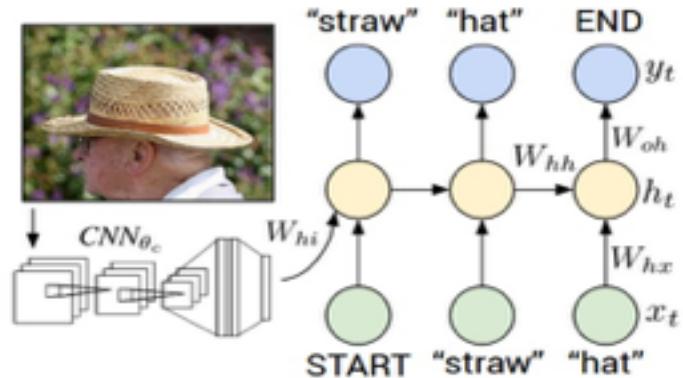
Amount of Data	very similar dataset	very different dataset
<b>very little data</b>	Use Linear Classifier on top layer	You're in trouble... Try linear classifier from different stages
<b>quite a lot of data</b>	Finetune a few layers	Finetune a larger number of layers

# Transfer learning with CNNs is pervasive... (it's the norm, not an exception)

Object  
Detection  
(Faster R-  
CNN)

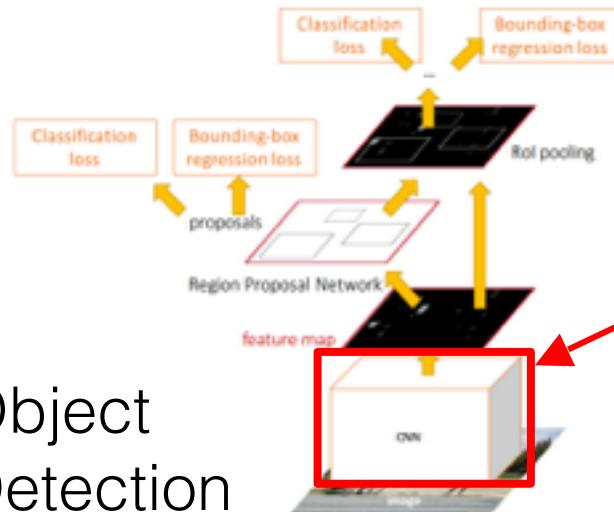


## Image Captioning: CNN + RNN



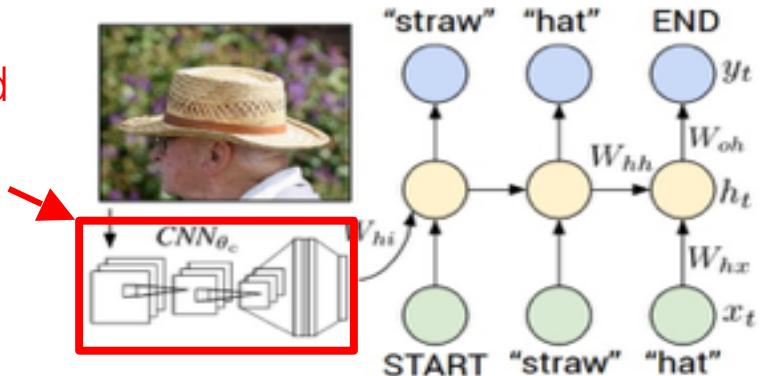
# Transfer learning with CNNs is pervasive... (it's the norm, not an exception)

Object  
Detection  
(Faster R-  
CNN)



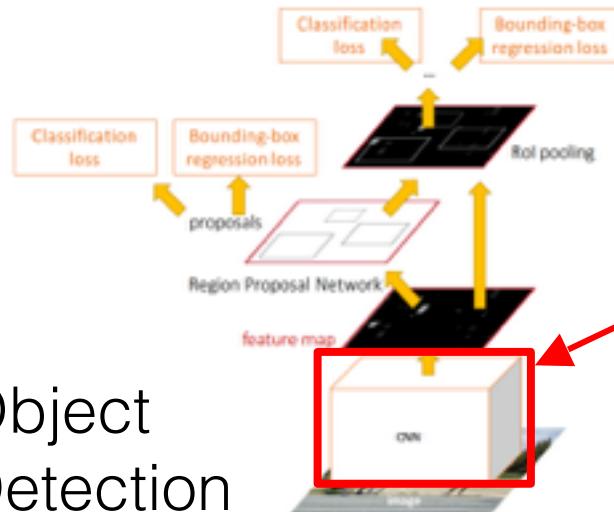
CNN pretrained  
on ImageNet

Image Captioning:  
CNN + RNN



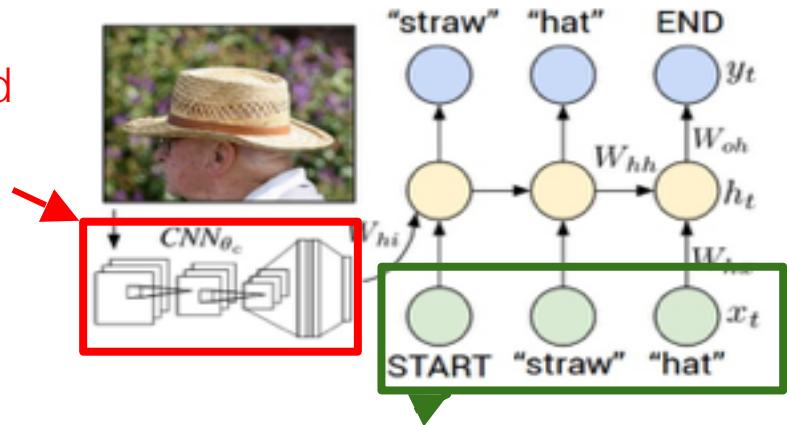
# Transfer learning with CNNs is pervasive... (it's the norm, not an exception)

Object  
Detection  
(Faster R-  
CNN)



CNN pretrained  
on ImageNet

Image Captioning:  
CNN + RNN



Word vectors  
pretrained from  
word2vec

# **Takeaway for your projects/beyond:**

Have some dataset of interest but it has  $< \sim 1M$  images?

1. Find a very large dataset that has similar data, train a big ConvNet there.
2. Transfer learn to your dataset

Caffe ConvNet library has a “**Model Zoo**” of pretrained models:

<https://github.com/BVLC/caffe/wiki/Model-Zoo>

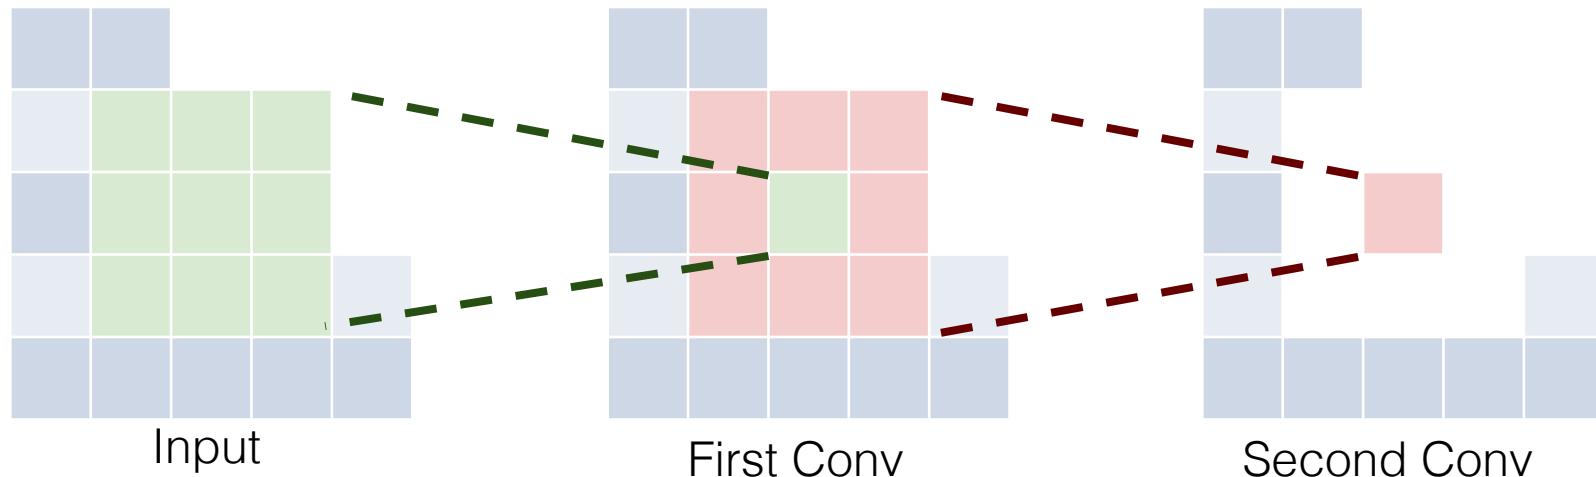
# All About Convolutions

## Part I: How to stack them

# The power of small filters

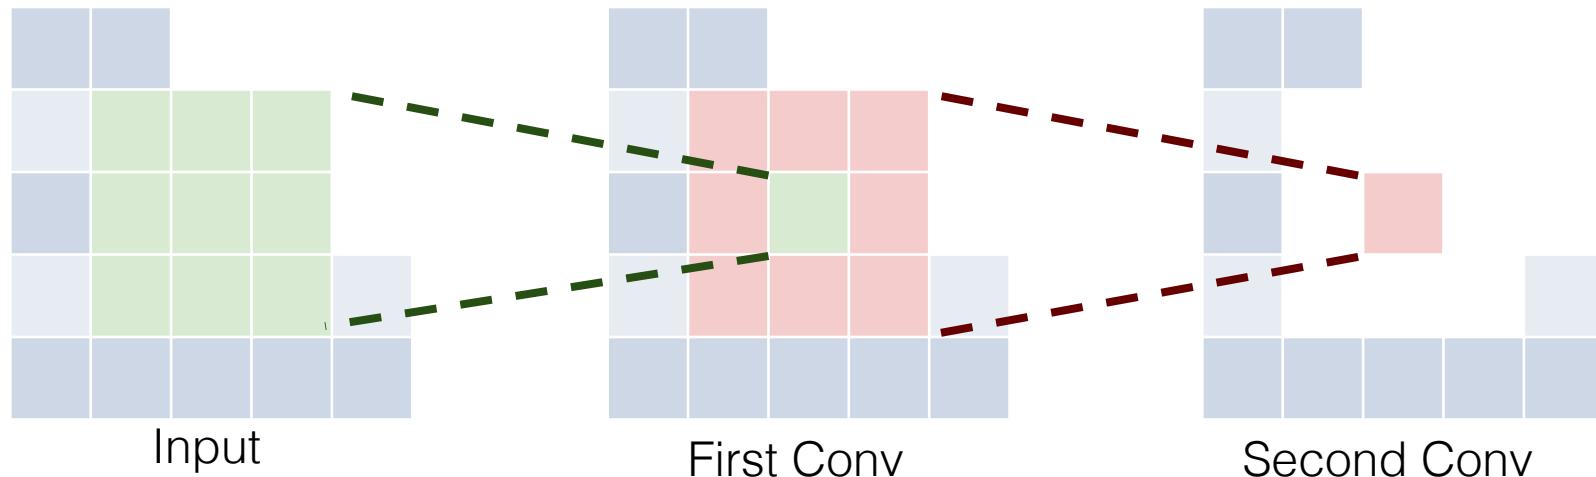
Suppose we stack two  $3 \times 3$  conv layers (stride 1)

Each neuron sees  $3 \times 3$  region of previous activation map



# The power of small filters

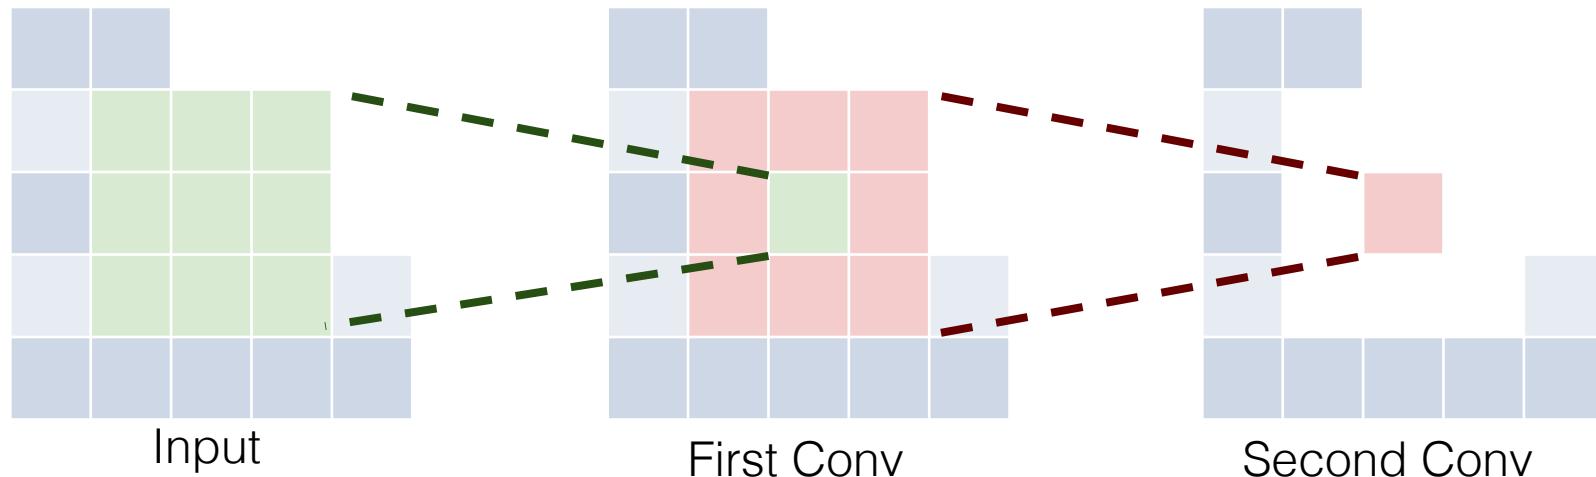
Question: How big of a region in the input does a neuron on the second conv layer see?



# The power of small filters

Question: How big of a region in the input does a neuron on the second conv layer see?

Answer:  $5 \times 5$



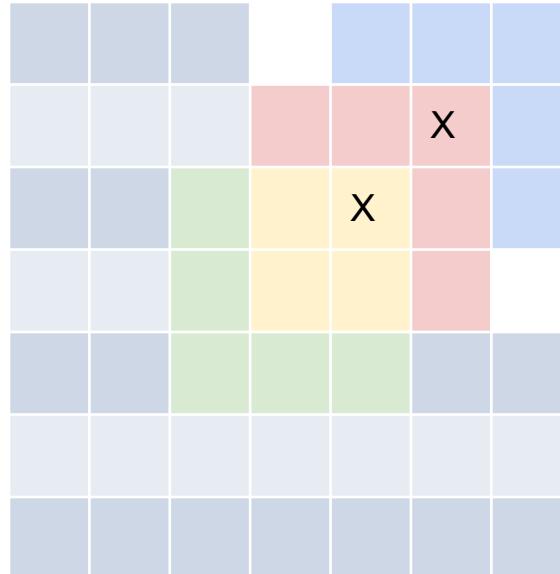
# The power of small filters

Question: If we stack three 3x3 conv layers, how big of an input region does a neuron in the third layer see?

# The power of small filters

Question: If we stack three 3x3 conv layers, how big of an input region does a neuron in the third layer see?

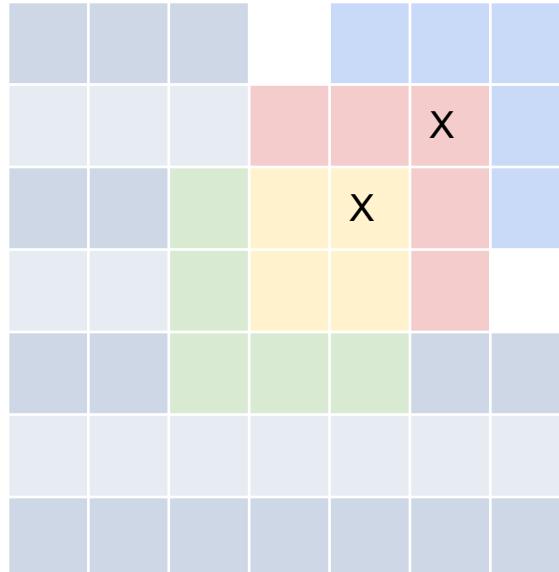
Answer:  $7 \times 7$



# The power of small filters

Question: If we stack three 3x3 conv layers, how big of an input region does a neuron in the third layer see?

Answer:  $7 \times 7$



Three  $3 \times 3$  conv  
gives similar  
representational  
power as a single  
 $7 \times 7$  convolution

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:

three CONV with  $3 \times 3$  filters

Number of weights:

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:

$$= C \times (7 \times 7 \times C) = \mathbf{49} C^2$$

three CONV with  $3 \times 3$  filters

Number of weights:

$$= 3 \times C \times (3 \times 3 \times C) = \mathbf{27} C^2$$

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:

$$= C \times (7 \times 7 \times C) = \mathbf{49 C^2}$$

three CONV with  $3 \times 3$  filters

Number of weights:

$$= 3 \times C \times (3 \times 3 \times C) = \mathbf{27 C^2}$$

Fewer parameters, more nonlinearity = GOOD

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:

$$= C \times (7 \times 7 \times C) = 49 C^2$$

Number of multiply-adds:

three CONV with  $3 \times 3$  filters

Number of weights:

$$= 3 \times C \times (3 \times 3 \times C) = 27 C^2$$

Number of multiply-adds:

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:

$$= C \times (7 \times 7 \times C) = 49 C^2$$

Number of multiply-adds:

$$\begin{aligned} &= (H \times W \times C) \times (7 \times 7 \times C) \\ &= \mathbf{49 HWC^2} \end{aligned}$$

three CONV with  $3 \times 3$  filters

Number of weights:

$$= 3 \times C \times (3 \times 3 \times C) = 27 C^2$$

Number of multiply-adds:

$$\begin{aligned} &= 3 \times (H \times W \times C) \times (3 \times 3 \times C) \\ &= \mathbf{27 HWC^2} \end{aligned}$$

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:

$$= C \times (7 \times 7 \times C) = 49 C^2$$

Number of multiply-adds:

$$= \mathbf{49 HWC^2}$$

three CONV with  $3 \times 3$  filters

Number of weights:

$$= 3 \times C \times (3 \times 3 \times C) = 27 C^2$$

Number of multiply-adds:

$$= \mathbf{27 HWC^2}$$



Less compute, more nonlinearity = GOOD

# The power of small filters

Why stop at  $3 \times 3$  filters? Why not try  $1 \times 1$ ?

# The power of small filters

Why stop at  $3 \times 3$  filters? Why not try  $1 \times 1$ ?

(note:  $1 \times 1$  filters sum across all channels of the input)

$$H \times W \times C$$

Conv 1x1, C/2 filters



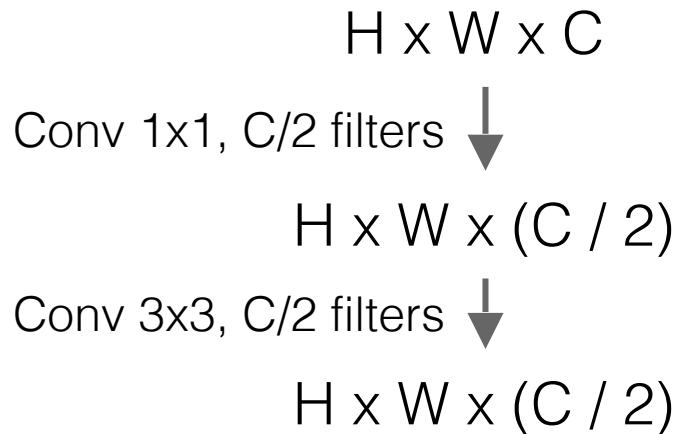
$$H \times W \times (C / 2)$$

1. “bottleneck”  $1 \times 1$  conv to reduce dimension

# The power of small filters

Why stop at  $3 \times 3$  filters? Why not try  $1 \times 1$ ?

(note:  $1 \times 1$  filters sum across all channels of the input)

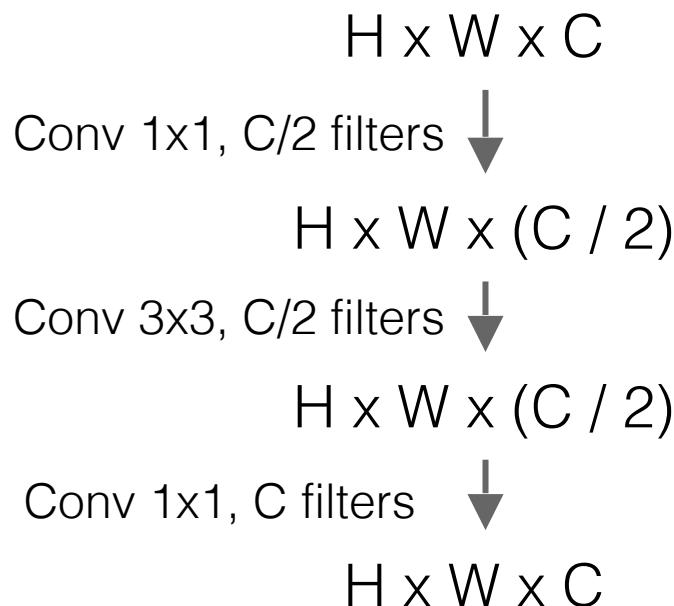


1. “bottleneck”  $1 \times 1$  conv to reduce dimension
2.  $3 \times 3$  conv at reduced dimension

# The power of small filters

Why stop at  $3 \times 3$  filters? Why not try  $1 \times 1$ ?

(note:  $1 \times 1$  filters sum across all channels of the input)

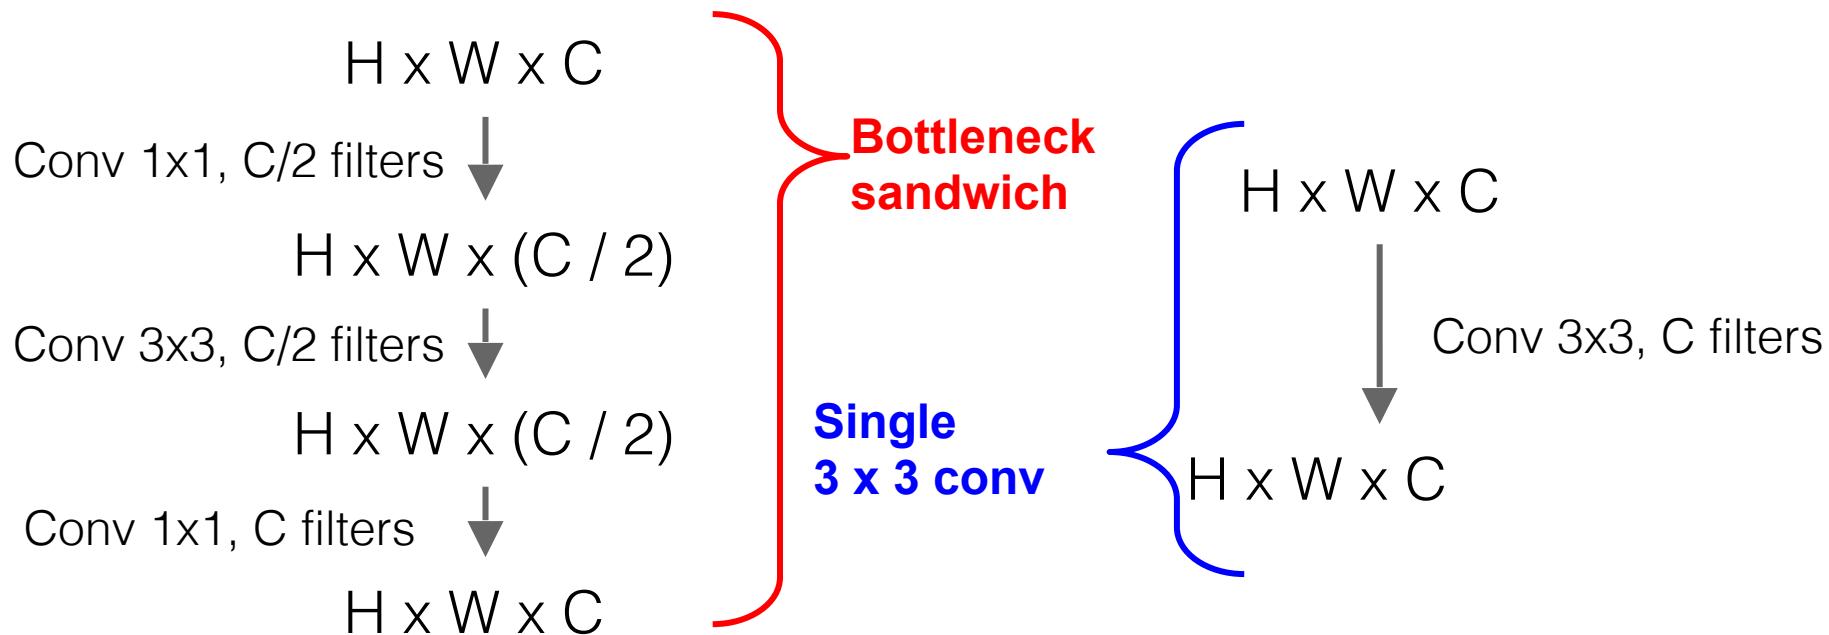


1. “bottleneck”  $1 \times 1$  conv to reduce dimension
2.  $3 \times 3$  conv at reduced dimension
3. Restore dimension with another  $1 \times 1$  conv

[Seen in Lin et al, “Network in Network”,  
GoogLeNet, ResNet]

# The power of small filters

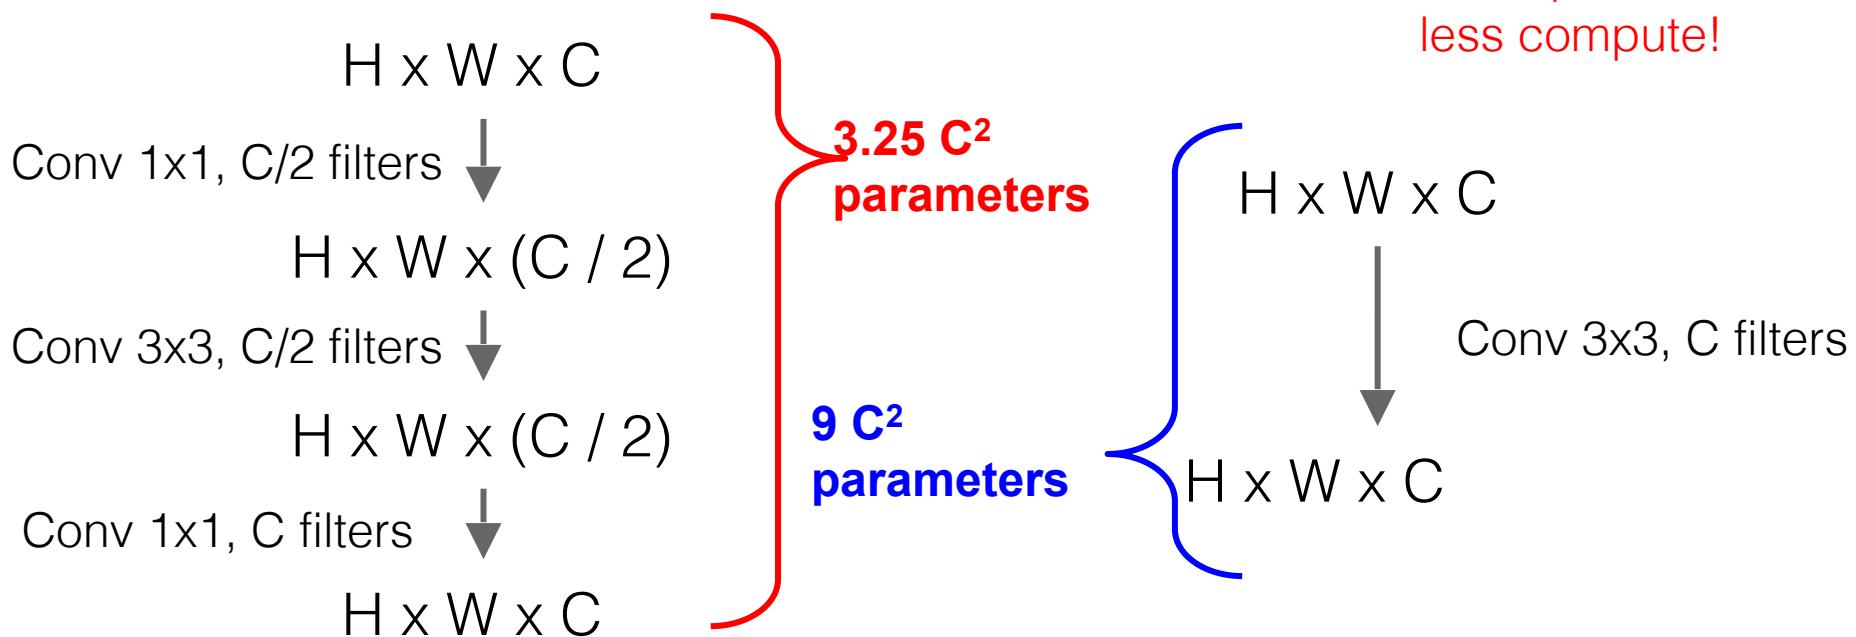
Why stop at  $3 \times 3$  filters? Why not try  $1 \times 1$ ?



# The power of small filters

Why stop at  $3 \times 3$  filters? Why not try  $1 \times 1$ ?

More nonlinearity,  
fewer params,  
less compute!

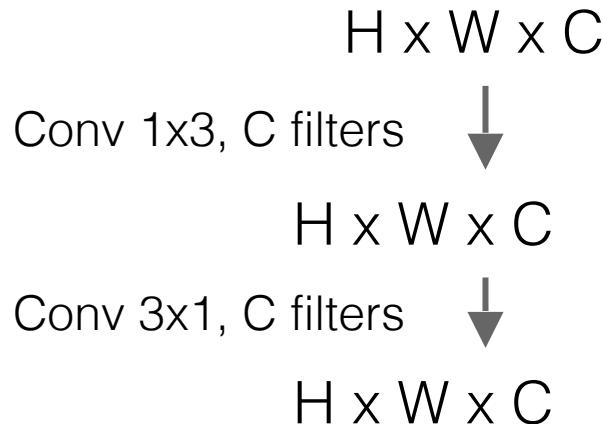


# The power of small filters

Still using 3 x 3 filters ... can we break it up?

# The power of small filters

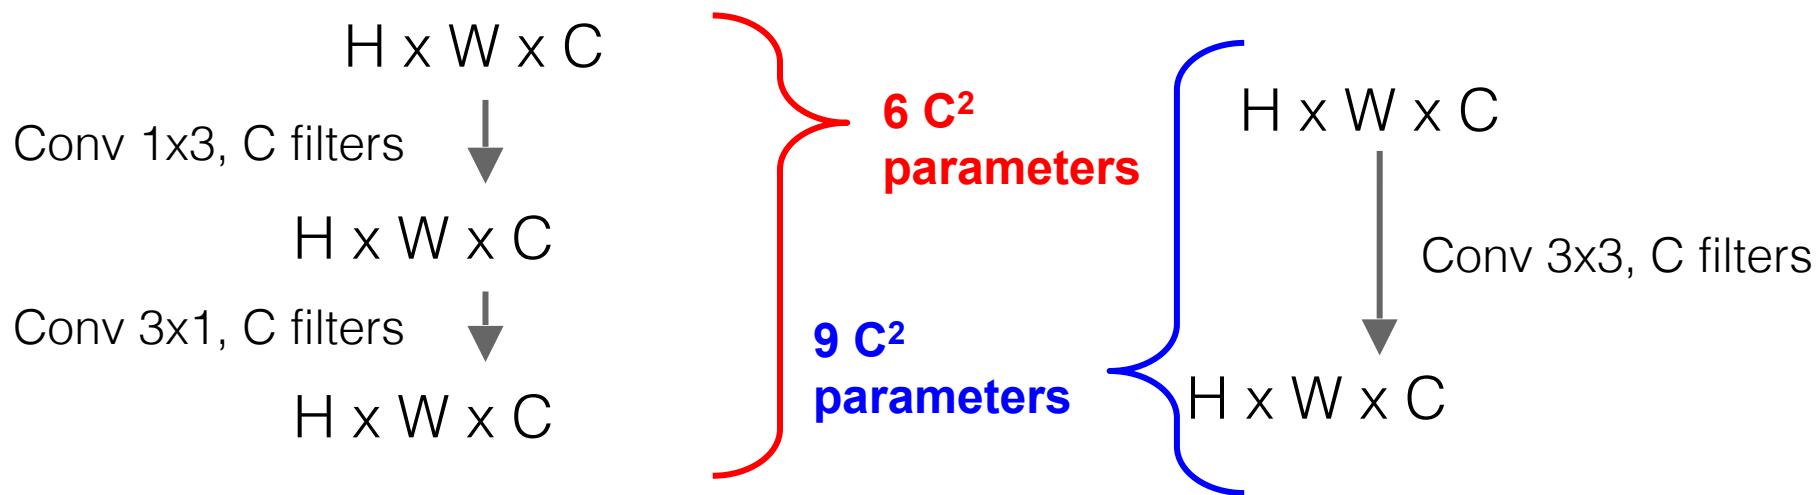
Still using  $3 \times 3$  filters ... can we break it up?



# The power of small filters

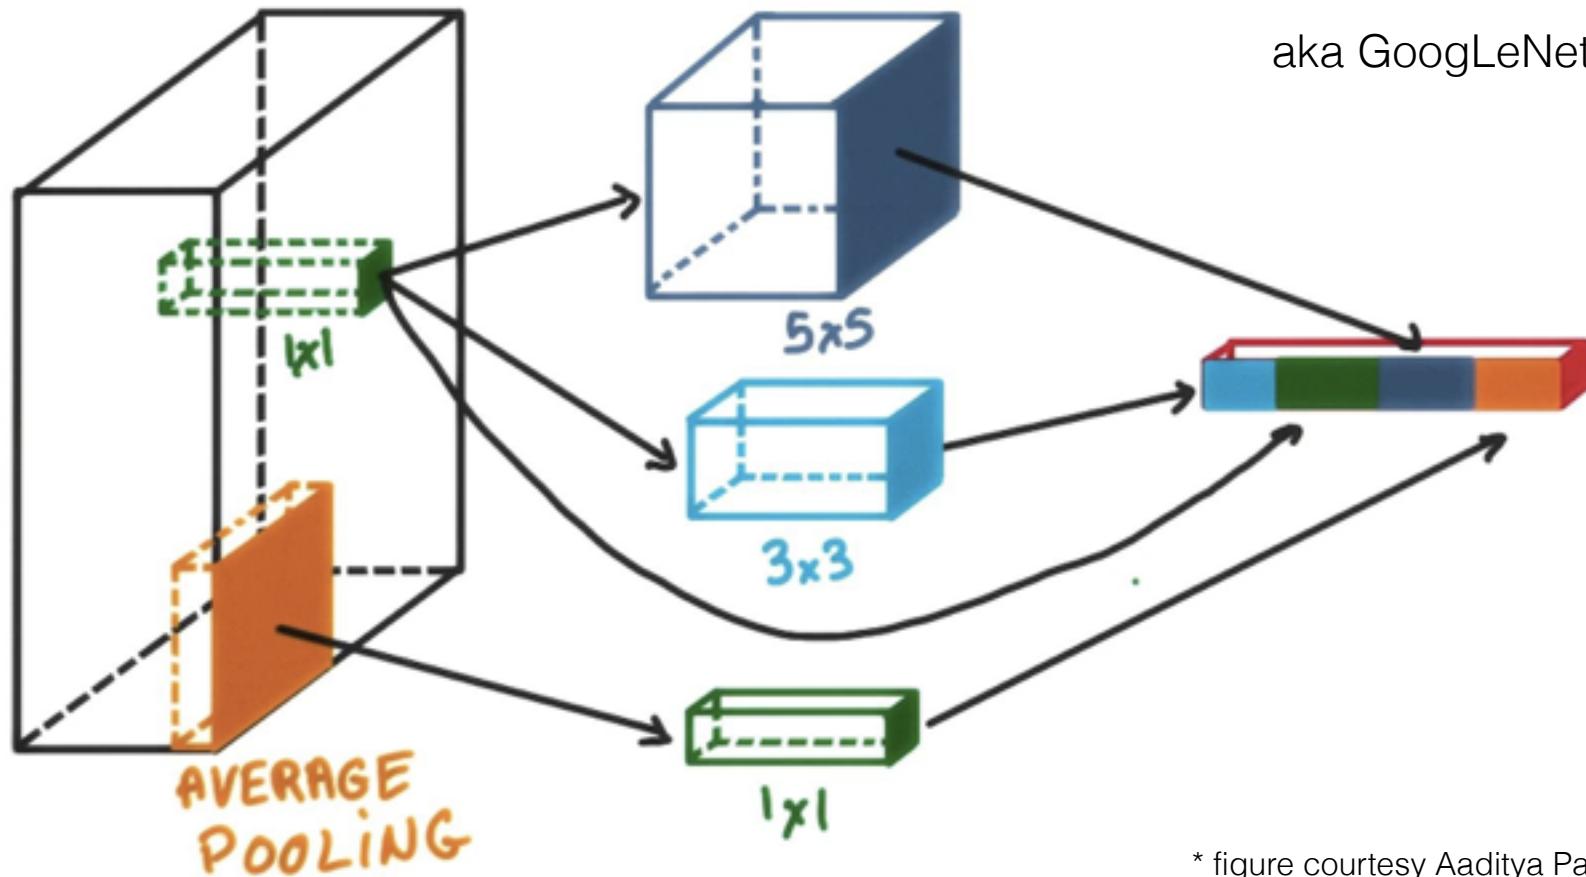
Still using  $3 \times 3$  filters ... can we break it up?

More nonlinearity,  
fewer params,  
less compute!



# INCEPTION MODULES

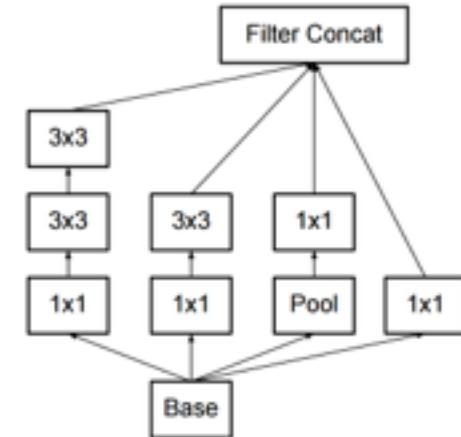
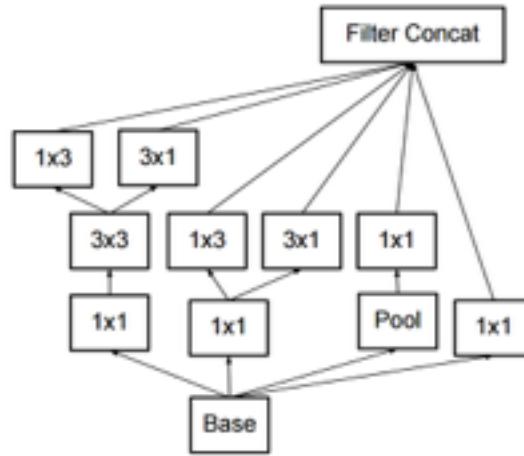
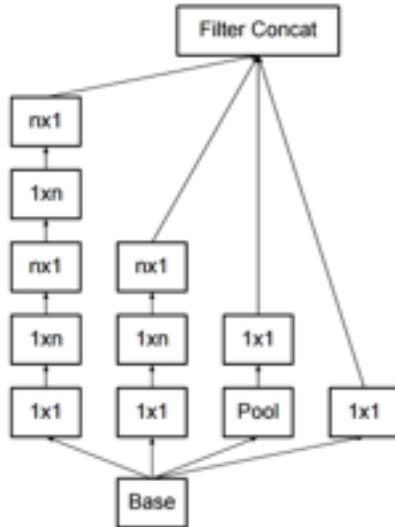
aka GoogLeNet



\* figure courtesy Aaditya Parkash

# The power of small filters

Latest version of GoogLeNet incorporates all these ideas



Szegedy et al, "Rethinking the Inception Architecture for Computer Vision"

# How to stack convolutions: Recap

- Replace large convolutions ( $5 \times 5$ ,  $7 \times 7$ ) with stacks of  $3 \times 3$  convolutions
- $1 \times 1$  “bottleneck” convolutions are very efficient
- Can factor  $N \times N$  convolutions into  $1 \times N$  and  $N \times 1$
- All of the above give fewer parameters, less compute, more nonlinearity

# All About Convolutions

## Part II: How to compute them

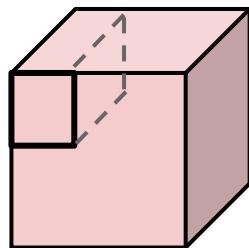
# Implementing Convolutions: im2col

There are highly optimized matrix multiplication routines for just about every platform

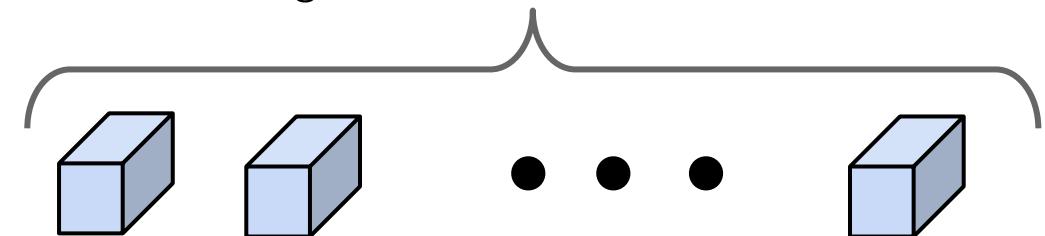
Can we turn convolution into matrix multiplication?

# Implementing Convolutions: im2col

Feature map:  $H \times W \times C$

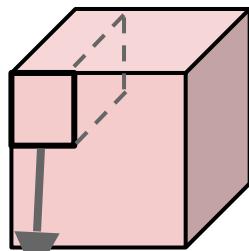


Conv weights: D filters, each  $K \times K \times C$



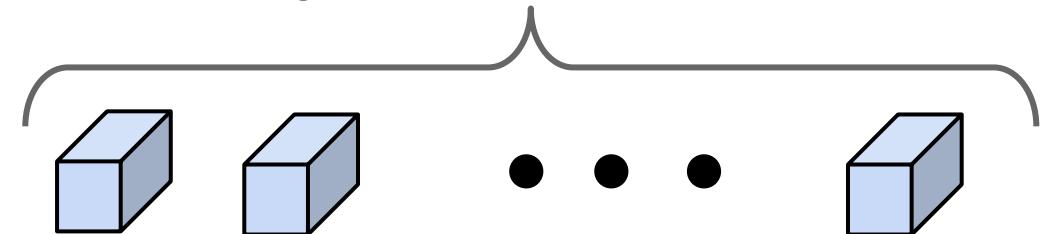
# Implementing Convolutions: im2col

Feature map:  $H \times W \times C$



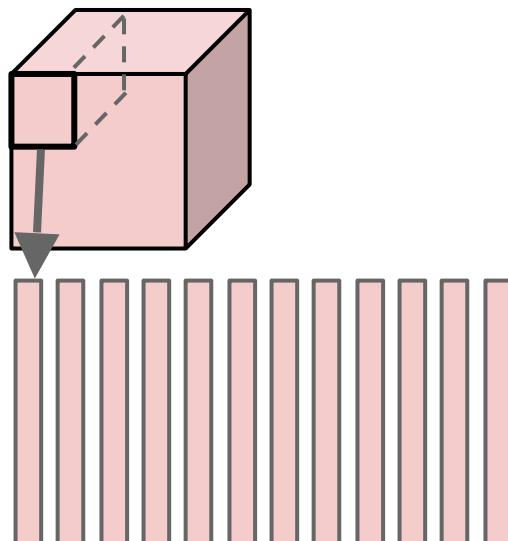
Reshape  $K \times K \times C$   
receptive field  
column with  $K^2C$   
elements

Conv weights:  $D$  filters, each  $K \times K \times C$

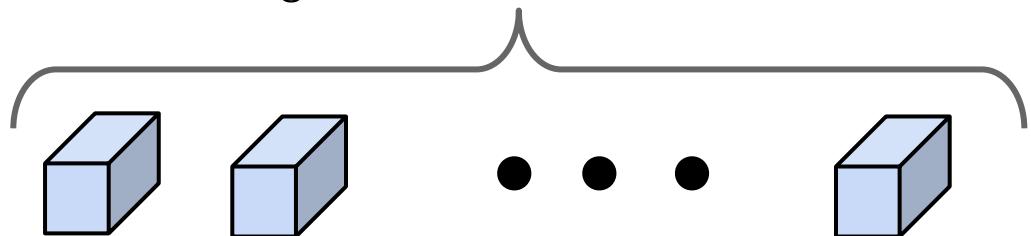


# Implementing Convolutions: im2col

Feature map:  $H \times W \times C$



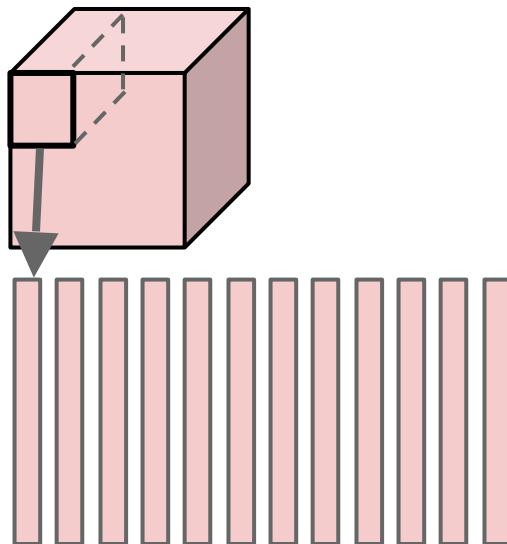
Conv weights: D filters, each  $K \times K \times C$



Repeat for all columns to get  $(K^2C) \times N$  matrix  
( $N$  receptive field locations)

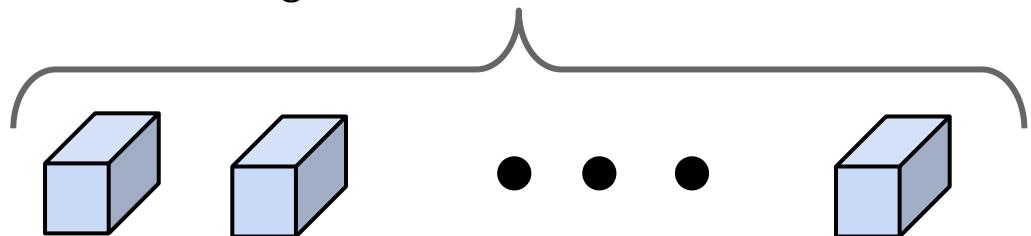
# Implementing Convolutions: im2col

Feature map:  $H \times W \times C$



Repeat for all columns to get  $(K^2C) \times N$  matrix  
( $N$  receptive field locations)

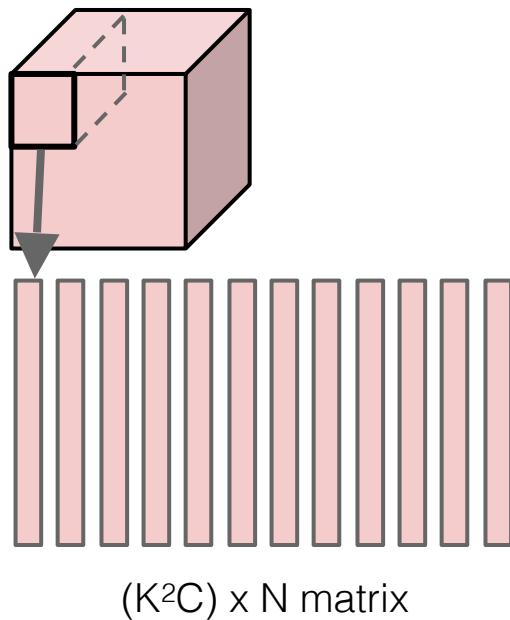
Conv weights:  $D$  filters, each  $K \times K \times C$



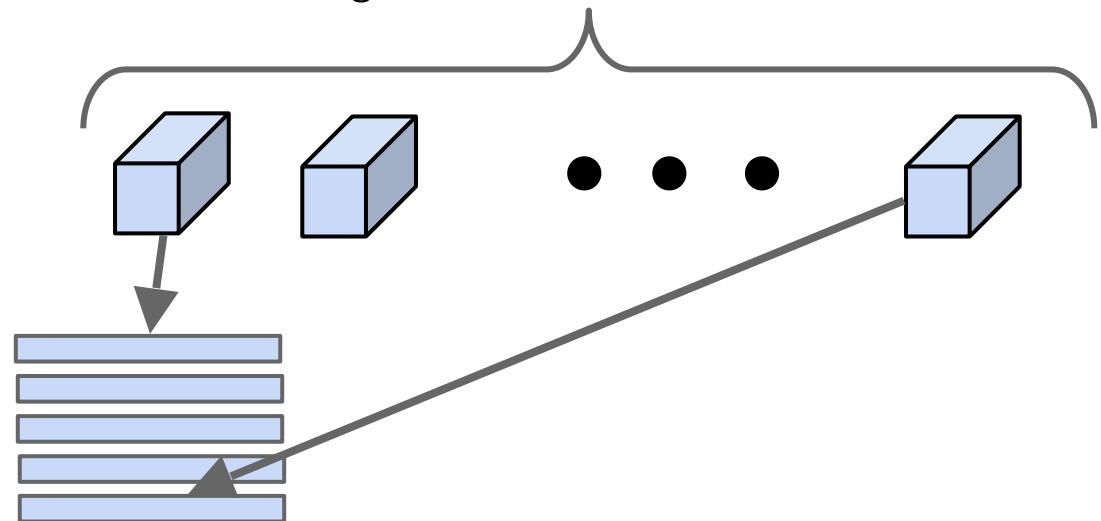
Elements appearing in  
multiple receptive fields are  
duplicated; this uses a lot of  
memory

# Implementing Convolutions: im2col

Feature map:  $H \times W \times C$

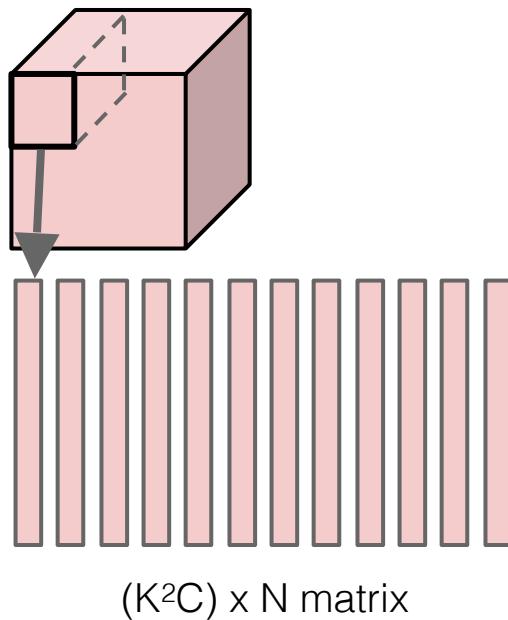


Conv weights: D filters, each  $K \times K \times C$

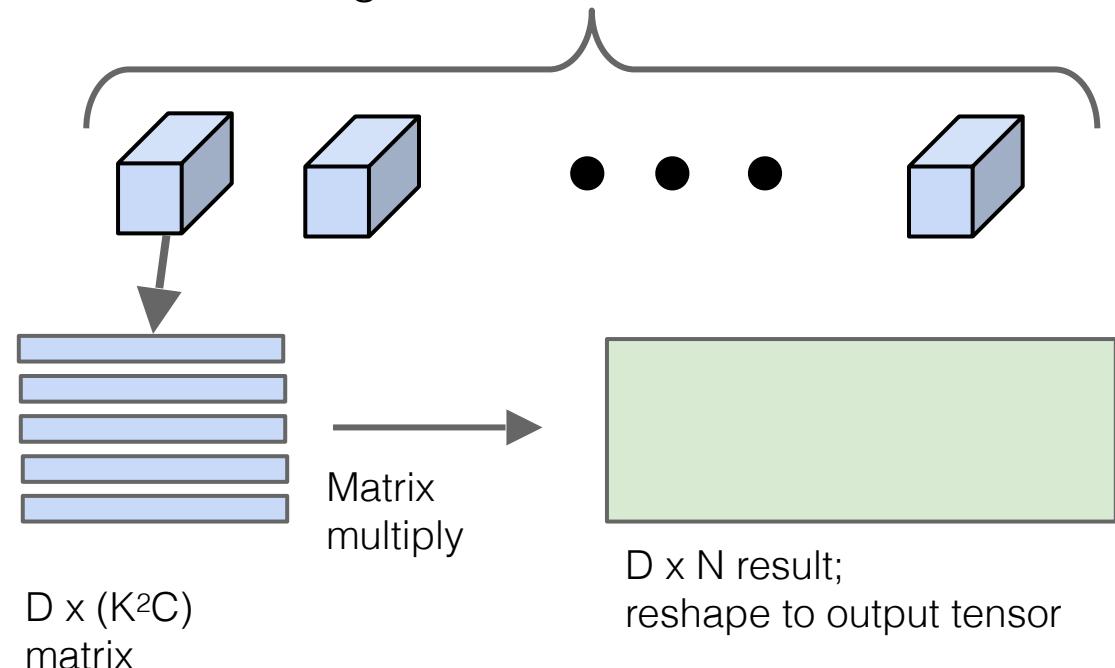


# Implementing Convolutions: im2col

Feature map:  $H \times W \times C$



Conv weights: D filters, each  $K \times K \times C$



```

template <typename Dtype>
void ConvolutionLayer<Dtype>::Forward_gpu(const vector<Blob<Dtype>*>& bottom,
    vector<Blob<Dtype>*>& top) {
  for (int i = 0; i < bottom.size(); ++i) {
    const Dtype* bottom_data = bottom[i]->gpu_data();
    Dtype* top_data = (*top)[i]->mutable_gpu_data();
    Dtype* col_data = col_buffer_.mutable_gpu_data();
    const Dtype* weight = this->blobs_[0]->gpu_data();
    int weight_offset = M_ * K_;
    int col_offset = K_ * N_;
    int top_offset = M_ * N_;
    for (int n = 0; n < num_; ++n) {
      // im2col transformation: unroll input regions for filtering
      // into column matrix for multiplication
      im2col_gpu(bottom_data + bottom[i]->offset(n), channels_, height_,
                 width_, kernel_h_, kernel_w_, pad_h_, pad_w_, stride_h_, stride_w_,
                 col_data);
      // Take inner products for groups.
      for (int g = 0; g < group_; ++g) {
        caffe_gpu_gemm<Dtype>(CblasNoTrans, CblasNoTrans, M_, N_, K_,
                               (Dtype)1., weight + weight_offset * g, col_data + col_offset * g,
                               (Dtype)0., top_data + (*top)[i]->offset(n) + top_offset * g);
      }
      // Add bias.
      if (bias_term_) {
        caffe_gpu_gemm<Dtype>(CblasNoTrans, CblasNoTrans, num_output_,
                               N_, 1, (Dtype)1., this->blobs_[1]->gpu_data(),
                               bias_multiplier_.gpu_data(),
                               (Dtype)1., top_data + (*top)[i]->offset(n));
      }
    }
  }
}

```

# Case study: CONV forward in Caffe library

im2col

matrix multiply: call to  
cuBLAS

bias offset

```

def conv_forward_strides(x, w, b, conv_param):
    N, C, H, W = x.shape
    F, HH, WW = w.shape
    stride, pad = conv_param['stride'], conv_param['pad']

    # Check dimensions
    assert (W + 2 * pad - WW) % stride == 0, 'width does not work'
    assert (H + 2 * pad - HH) % stride == 0, 'height does not work'

    # Pad the input
    p = pad
    x_padded = np.pad(x, ((0, 0), (0, 0), (p, p), (p, p)), mode='constant')

    # Figure out output dimensions
    H += 2 * pad
    W += 2 * pad
    out_h = (H - HH) / stride + 1
    out_w = (W - WW) / stride + 1

    # Perform an im2col operation by picking clever strides
    shape = (C, HH, WW, N, out_h, out_w)
    strides = (W * H, W, 1, C * H * W, stride * W, stride)
    strides = x.itemsize * np.array(strides)
    x_stride = np.lib.stride_tricks.as_strided(x_padded,
                                                shape=shape, strides=strides)
    x_cols = np.ascontiguousarray(x_stride)
    x_cols.shape = (C * HH * WW, N * out_h * out_w)

    # Now all our convolutions are a big matrix multiply
    res = w.reshape(F, -1).dot(x_cols) + b.reshape(-1, 1)

    # Reshape the output
    res.shape = (F, N, out_h, out_w)
    out = res.transpose(1, 0, 2, 3)

    # Be nice and return a contiguous array
    # The old version of conv_forward_fast doesn't do this, so for a fair
    # comparison we won't either
    out = np.ascontiguousarray(out)

    cache = (x, w, b, conv_param, x_cols)
    return out, cache

```

# Case study: fast\_layers.py from HW

im2col

matrix multiply:  
call np.dot  
(which calls BLAS)

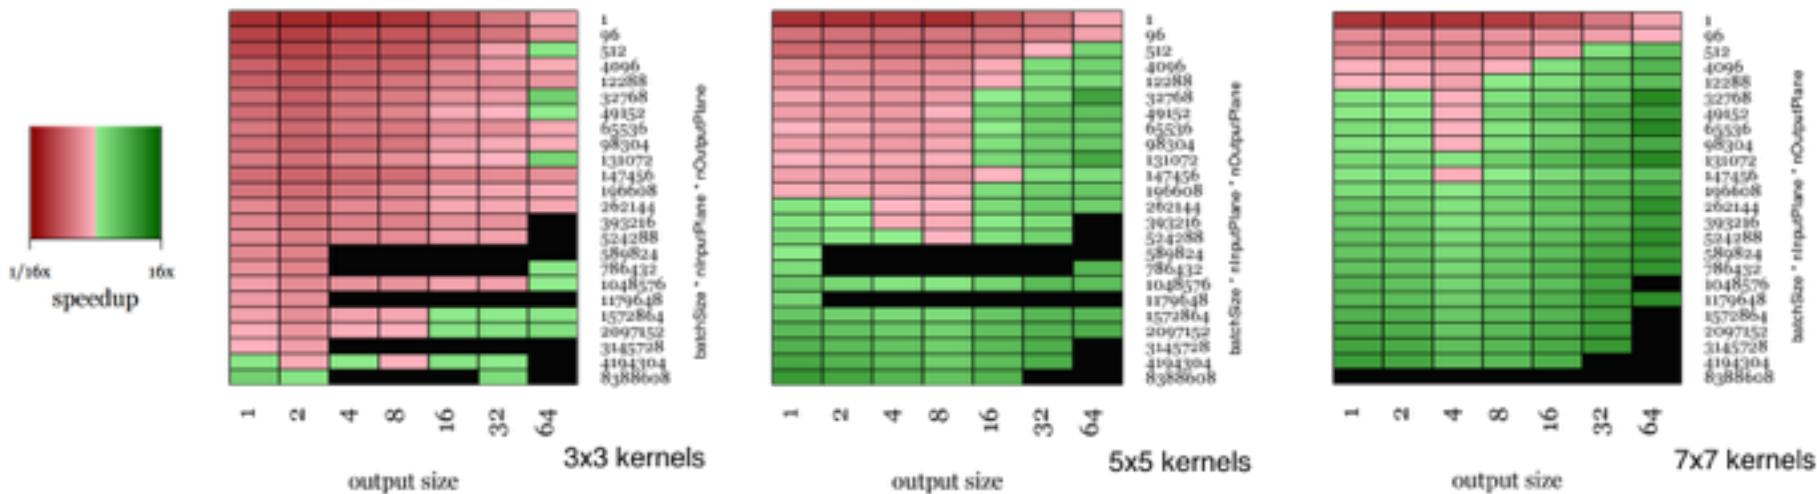
# Implementing convolutions: FFT

- **Convolution Theorem:** The convolution of  $f$  and  $g$  is equal to the elementwise product of their Fourier Transforms:  $\mathcal{F}(f * g) = \mathcal{F}(f) \cdot \mathcal{F}(g)$
- Using the **Fast Fourier Transform**, we can compute the Discrete Fourier transform of an  $N$ -dimensional vector in  $O(N \log N)$  time (also extends to 2D images)

# Implementing convolutions: FFT

1. Compute FFT of weights:  $F(W)$
2. Compute FFT of image:  $F(X)$
3. Compute elementwise product:  $F(W) \circ F(X)$
4. Compute inverse FFT:  $Y = F^{-1}(F(W) \circ F(X))$

# Implementing convolutions: FFT



FFT convolutions get a big speedup for larger filters  
Not much speedup for 3x3 filters =(

Vasilache et al, Fast Convolutional Nets With fbfft: A GPU Performance Evaluation

\* Original slides borrowed from Andrej Karpathy  
and Li Fei-Fei, Stanford cs231n

# Implementing convolution: “Fast Algorithms”

**Naive matrix multiplication:** Computing product of two  $N \times N$  matrices takes  $O(N^3)$  operations

**Strassen’s Algorithm:** Use clever arithmetic to reduce complexity to  $O(N^{\log 2(7)}) \sim O(N^{2.81})$

$$A = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix}$$

$$M_1 := (A_{1,1} + A_{2,2})(B_{1,1} + B_{2,2})$$

$$B = \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix}$$

$$M_2 := (A_{2,1} + A_{2,2})B_{1,1}$$

$$M_3 := A_{1,1}(B_{1,2} - B_{2,2})$$

$$M_4 := A_{2,2}(B_{2,1} - B_{1,1})$$

$$C = \begin{bmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{bmatrix}$$

$$M_5 := (A_{1,1} + A_{1,2})B_{2,2}$$

$$M_6 := (A_{2,1} - A_{1,1})(B_{1,1} + B_{1,2})$$

$$M_7 := (A_{1,2} - A_{2,2})(B_{2,1} + B_{2,2})$$

$$C_{1,1} = M_1 + M_4 - M_5 + M_7$$

$$C_{1,2} = M_3 + M_5$$

$$C_{2,1} = M_2 + M_4$$

$$C_{2,2} = M_1 - M_2 + M_3 + M_6$$

From Wikipedia

# Implementing convolution: “Fast Algorithms”

Similar cleverness can be applied to convolutions

Lavin and Gray (2015) work out special cases for 3x3 convolutions:

$$F(2, 3) = \begin{bmatrix} d_0 & d_1 & d_2 \\ d_1 & d_2 & d_3 \end{bmatrix} \begin{bmatrix} g_0 \\ g_1 \\ g_2 \end{bmatrix} = \begin{bmatrix} m_1 + m_2 + m_3 \\ m_2 - m_3 - m_4 \end{bmatrix}$$

$$\begin{aligned} m_1 &= (d_0 - d_2)g_0 & m_2 &= (d_1 + d_2)\frac{g_0 + g_1 + g_2}{2} \\ m_4 &= (d_1 - d_3)g_2 & m_3 &= (d_2 - d_1)\frac{g_0 - g_1 + g_2}{2} \end{aligned}$$

$$B^T = \begin{bmatrix} 1 & 0 & -1 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & -1 & 1 & 0 \\ 0 & 1 & 0 & -1 \end{bmatrix}$$

$$G = \begin{bmatrix} 1 & 0 & 0 \\ \frac{1}{2} & \frac{1}{2} & \frac{1}{2} \\ \frac{1}{2} & -\frac{1}{2} & \frac{1}{2} \\ 0 & 0 & 1 \end{bmatrix}$$

$$A^T = \begin{bmatrix} 1 & 1 & 1 & 0 \\ 0 & 1 & -1 & -1 \end{bmatrix}$$

$$g = [g_0 \ g_1 \ g_2]^T$$

$$d = [d_0 \ d_1 \ d_2 \ d_3]^T$$

Lavin and Gray, “Fast Algorithms for Convolutional Neural Networks”, 2015

# Implementing convolution: “Fast Algorithms”

Huge speedups on VGG for small batches:

N	cuDNN		F( $2 \times 2, 3 \times 3$ )		Speedup
	msec	TFLOPS	msec	TFLOPS	
1	12.52	3.12	5.55	7.03	2.26X
2	20.36	3.83	9.89	7.89	2.06X
4	104.70	1.49	17.72	8.81	5.91X
8	241.21	1.29	33.11	9.43	7.28X
16	203.09	3.07	65.79	9.49	3.09X
32	237.05	5.27	132.36	9.43	1.79X
64	394.05	6.34	266.48	9.37	1.48X

Table 5. cuDNN versus  $F(2 \times 2, 3 \times 3)$  performance on VGG Network E with fp32 data. Throughput is measured in Effective TFLOPS, the ratio of direct algorithm GFLOPs to run time.

N	cuDNN		F( $2 \times 2, 3 \times 3$ )		Speedup
	msec	TFLOPS	msec	TFLOPS	
1	14.58	2.68	5.53	7.06	2.64X
2	20.94	3.73	9.83	7.94	2.13X
4	104.19	1.50	17.50	8.92	5.95X
8	241.87	1.29	32.61	9.57	7.42X
16	204.01	3.06	62.93	9.92	3.24X
32	236.13	5.29	123.12	10.14	1.92X
64	395.93	6.31	242.98	10.28	1.63X

Table 6. cuDNN versus  $F(2 \times 2, 3 \times 3)$  performance on VGG Network E with fp16 data.

# Computing Convolutions: Recap

- im2col: Easy to implement, but big memory overhead
- FFT: Big speedups for small kernels
- “Fast Algorithms” seem promising, not widely used yet

# Implementation Details



\* Original slides borrowed from Andrej Karpathy  
and Li Fei-Fei, Stanford cs231n

# Spot the CPU!



\* Original slides borrowed from Andrej Karpathy and Li Fei-Fei, Stanford cs231n

# Spot the CPU!

“central processing unit”



# Spot the GPU!

“graphics processing  
unit”



# Spot the GPU!

“graphics processing unit”





VS





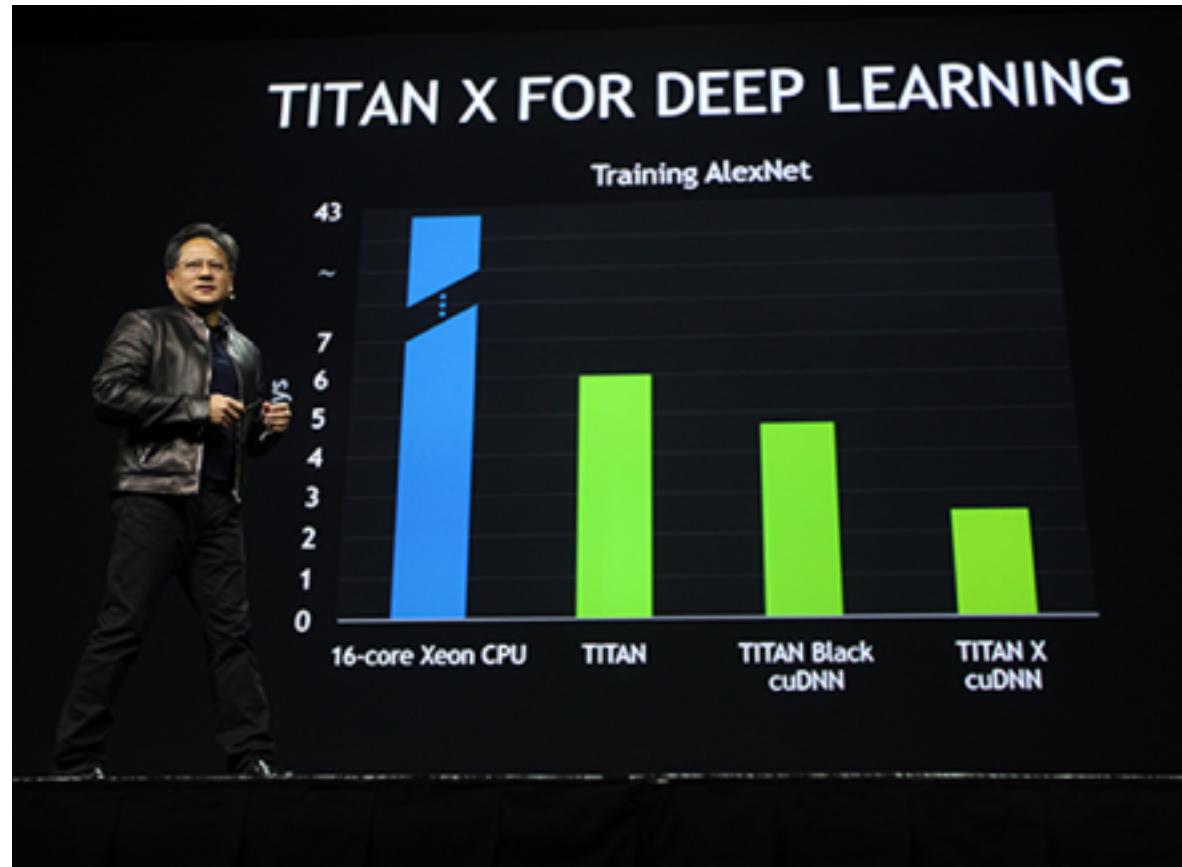
VS



NVIDIA is much more  
common for deep learning

## GTC 2015:

Introduced new Titan X GPU by bragging about AlexNet benchmarks



\* Original slides borrowed from Andrej Karpathy and Li Fei-Fei, Stanford cs231n

# CPU

Few, fast cores (1 - 16)

Good at sequential processing



# GPU

Many, slower cores (thousands)

Originally for graphics

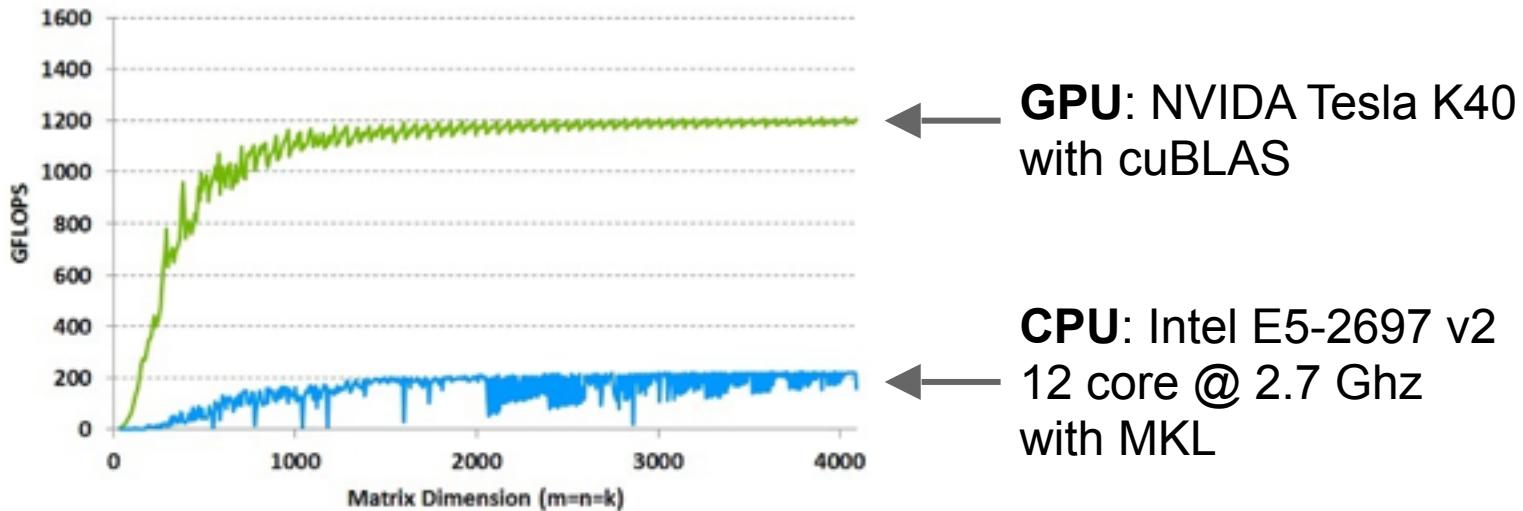
Good at parallel computation



# GPUs can be programmed

- CUDA (NVIDIA only)
  - Write C code that runs directly on the GPU
  - Higher-level APIs: cuBLAS, cuFFT, cuDNN, etc
- OpenCL
  - Similar to CUDA, but runs on anything
  - Usually slower :(
- Udacity: Intro to Parallel Programming [https://www.udacity.com/  
course/cs344](https://www.udacity.com/course/cs344)
  - For deep learning just use existing libraries

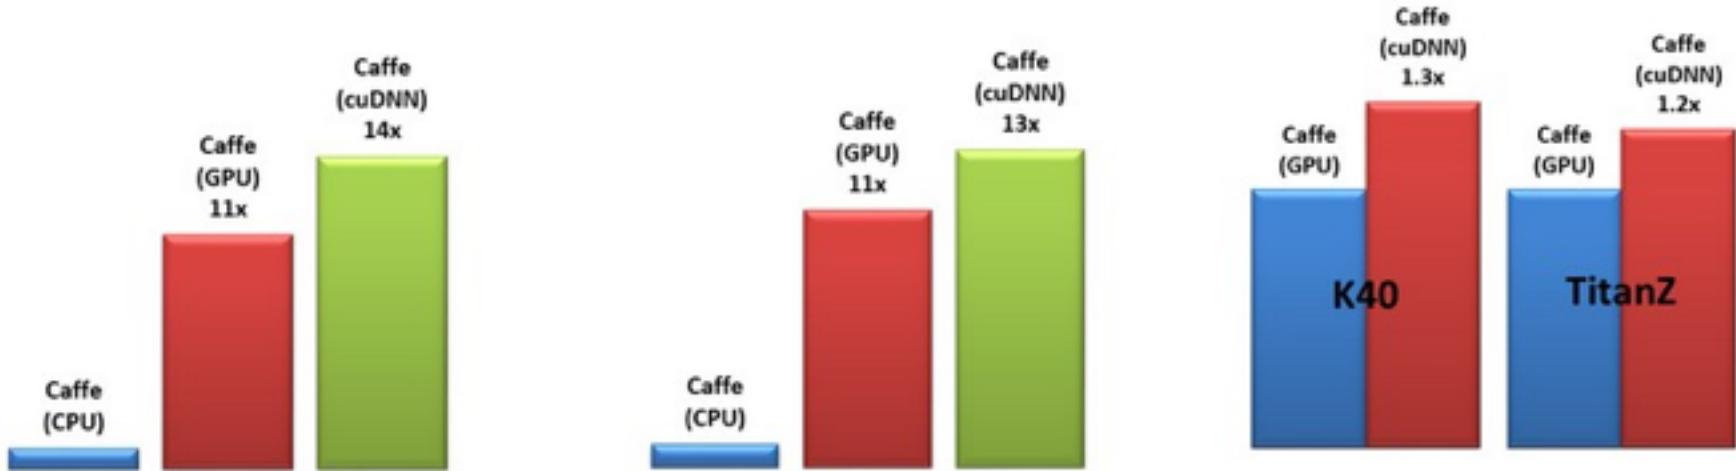
GPUs are really good  
at matrix multiplication:



**GPU:** NVIDIA Tesla K40  
with cuBLAS

**CPU:** Intel E5-2697 v2  
12 core @ 2.7 Ghz  
with MKL

GPUs are really good at convolution (cuDNN):



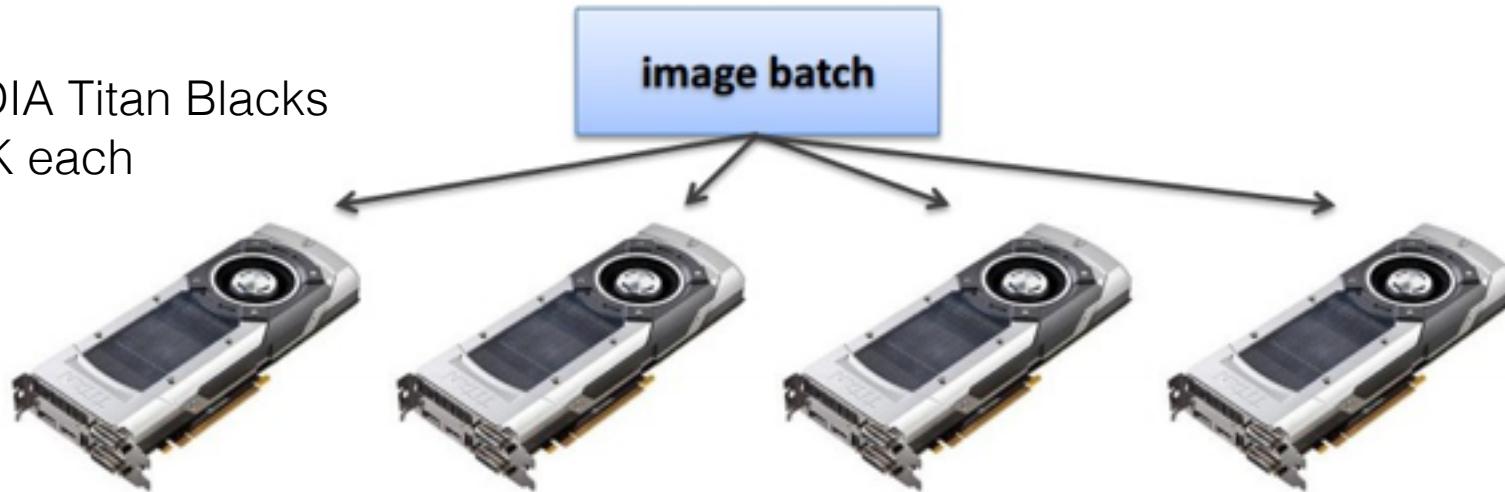
All comparisons are against a 12-core Intel E5-2679v2 CPU @ 2.4GHz running Caffe with Intel MKL 11.1.3.

# Even with GPUs, training can be slow

**VGG:** ~2-3 weeks training with 4 GPUs

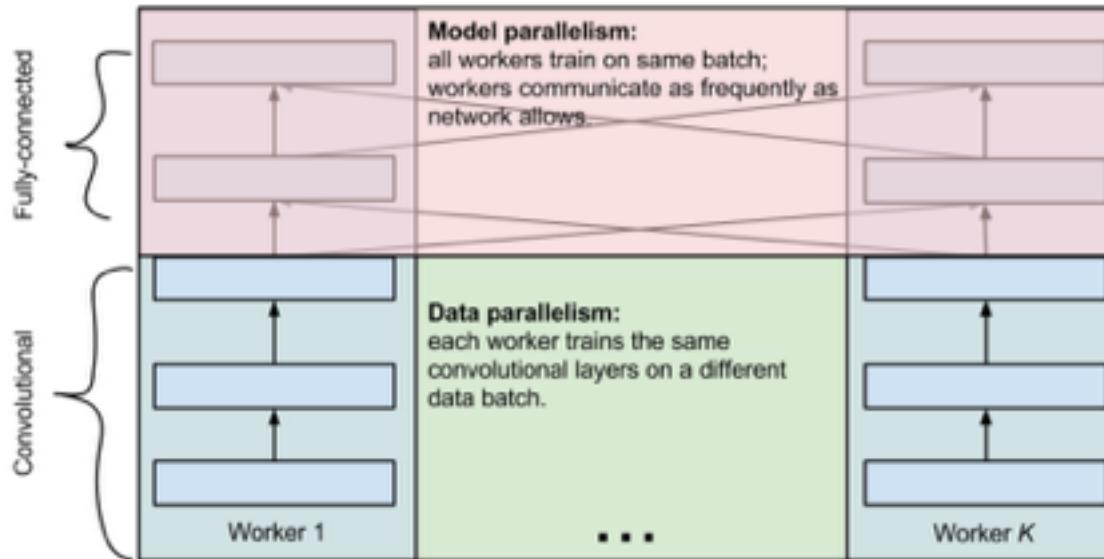
**ResNet 101:** 2-3 weeks with 4 GPUs

NVIDIA Titan Blacks  
~\$1K each



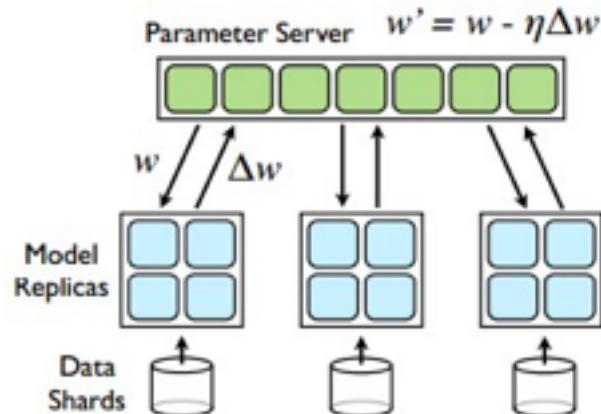
ResNet reimplemented in Torch: <http://torch.ch/blog/2016/02/04/resnets.html>

# Multi-GPU training: More complex



Alex Krizhevsky, “One weird trick for parallelizing convolutional neural networks”

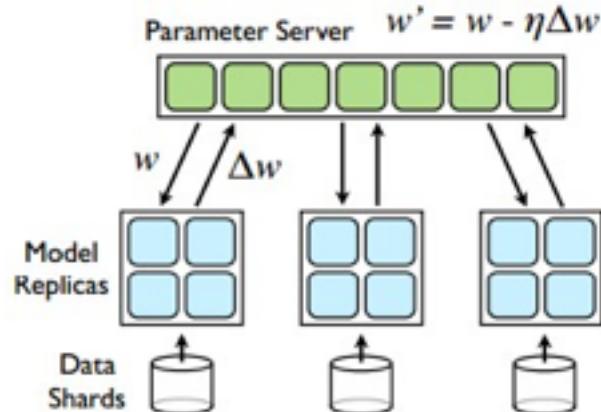
# Google: Distributed CPU training



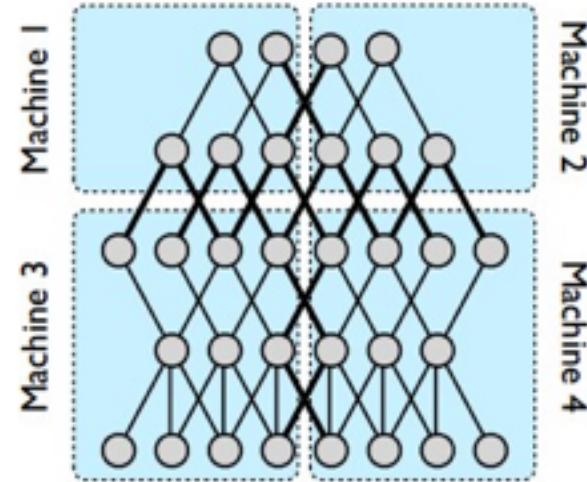
**Data parallelism**

*[Large Scale Distributed Deep Networks, Jeff Dean et al., 2013]*

# Google: Distributed CPU training



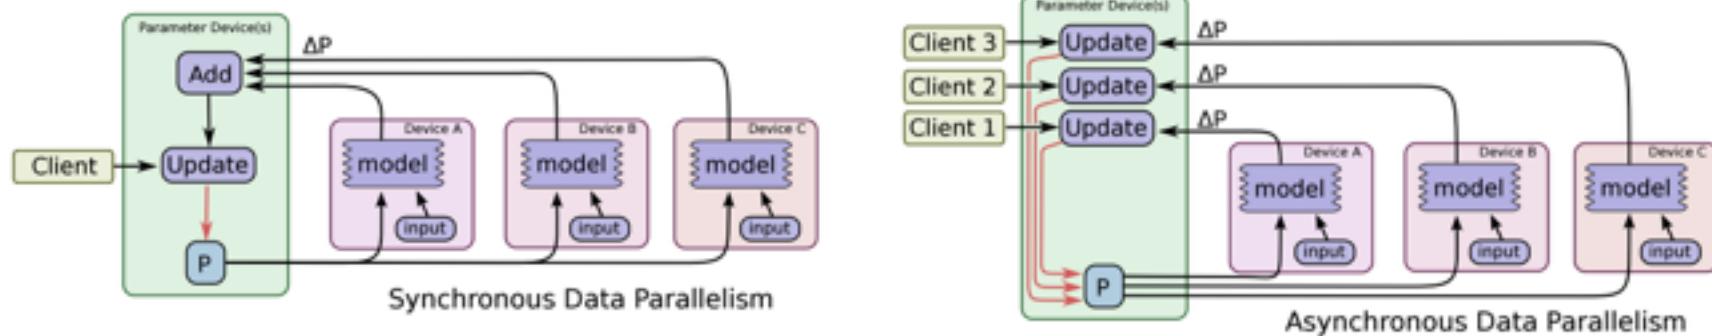
**Data parallelism**



**Model parallelism**

*[Large Scale Distributed Deep Networks, Jeff Dean et al., 2013]*

# Google: Synchronous vs Async



*Abadi et al, “TensorFlow: Large-Scale Machine Learning on Heterogeneous Distributed Systems”*

\* Original slides borrowed from Andrej Karpathy  
and Li Fei-Fei, Stanford cs231n

# Bottlenecks

to be aware of



# **GPU - CPU communication is a bottleneck.**

=>

**CPU** data prefetch+augment thread running

while

**GPU** performs forward/backward pass

# CPU - disk bottleneck

Hard disk is slow to read from

=> Pre-processed images  
stored contiguously in files, read as  
raw byte stream from SSD disk

Moving parts lol



# GPU memory bottleneck

Titan X: 12 GB <- currently the max  
GTX 980 Ti: 6 GB

e.g.

AlexNet: ~3GB needed with batch size 256

# Floating Point Precision

# Floating point precision

- 64 bit “double” precision is default in a lot of programming
- 32 bit “single” precision is typically used for CNNs for performance

# Floating point precision

- 64 bit “double” precision is default in a lot of programming
- 32 bit “single” precision is typically used for CNNs for performance
  - Including in your homework!

```
class FullyConnectedNet(object):  
    """  
    A fully-connected neural network with an arbitrary number of hidden layers,  
    ReLU nonlinearities, and a softmax loss function. This will also implement  
    dropout and batch normalization as options. For a network with L layers,  
    the architecture will be  
  
    [affine - [batch norm] - relu - [dropout]]  $\times$  (L - i) - affine - softmax  
  
    where batch normalization and dropout are optional, and the (...) block is  
    repeated L - i times.  
  
    Similar to the TwoLayerNet above, learnable parameters are stored in the  
    self.params dictionary and will be learned using the Solver class.  
    """  
  
    def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,  
                 dropout=0, use_batch_norm=False, reg=0.0,  
                 weight_scale=1e-2, dtype=np.float32, seed=None):
```

# Floating point precision

- **Prediction:** 16 bit “half” precision will be the new standard
  - Already supported in cuDNN
  - Nervana fp16 kernels are the fastest right now
  - Hardware support in next-gen NVIDIA cards (Pascal)
  - Not yet supported in *Torch*

AlexNet (One Weird Trick paper) - Input 128x3x224x224

Library	Class	Time (ms)	forward (ms)	backward (ms)
Nervana-fp16	ConvLayer	92	29	62
CuDNN[R3]-fp16 (Torch)	cudnn.SpatialConvolution	96	30	66
CuDNN[R3]-fp32 (Torch)	cudnn.SpatialConvolution	96	32	64

OxfordNet [Model-A] - Input 64x3x224x224

Library	Class	Time (ms)	forward (ms)	backward (ms)
Nervana-fp16	ConvLayer	529	167	362
Nervana-fp32	ConvLayer	590	180	410
CuDNN[R3]-fp16 (Torch)	cudnn.SpatialConvolution	615	179	436

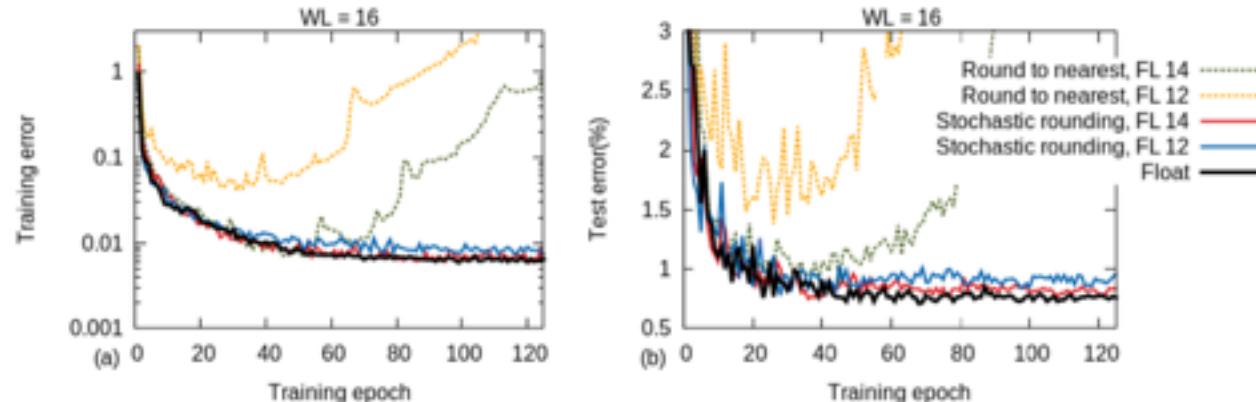
GoogleNet V1 - Input 128x3x224x224

Library	Class	Time (ms)	forward (ms)	backward (ms)
Nervana-fp16	ConvLayer	283	85	197
Nervana-fp32	ConvLayer	322	90	232
CuDNN[R3]-fp32 (Torch)	cudnn.SpatialConvolution	431	117	313

# Floating point precision

- How low can we go?
- Gupta et al, 2015:  
Train with **16-bit fixed point** with stochastic rounding

CNNs on  
MNIST



Gupta et al, "Deep Learning with Limited Numerical Precision", ICML 2015

# Floating point precision

- How low can we go?
- Courbariaux et al, 2015:  
Train with **10-bit activations, 12-bit parameter updates**

Courbariaux et al, "Training Deep Neural Networks with Low Precision Multiplications", ICLR 2015

# Floating point precision

- How low can we go?
- Courbariaux and Bengio, February 9 2016:
  - Train with **1-bit activations and weights!**
  - All activations and weights are +1 or -1
  - Fast multiplication with bitwise XNOR
  - (Gradients use higher precision)

Courbariaux et al, “BinaryNet: Training Deep Neural Networks with Weights and Activations Constrained to +1 or -1”, arXiv 2016

# Implementation details: Recap

- GPUs much faster than CPUs
- Distributed training is sometimes used
  - Not needed for small problems
- Be aware of bottlenecks: CPU / GPU, CPU / disk
- Low precision makes things faster and still works
  - 32 bit is standard now, 16 bit soon

# Recap

- Data augmentation: artificially expand your data
- Transfer learning: CNNs without huge data
- All about convolutions
- Implementation details

# TensorFlow

<https://www.tensorflow.org>

# TensorFlow

- From Google
- Very similar to Theano - all about computation graphs
- Easy visualizations (TensorBoard)
- Multi-GPU and multi-node training

# TensorFlow: Two-Layer Net

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D])
7 y = tf.placeholder(tf.float32, shape=[None, C])
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32))
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32))
11
12 a = tf.matmul(x, w1)
13 a_relu = tf.nn.relu(a)
14 scores = tf.matmul(a_relu, w2)
15 probs = tf.nn.softmax(scores)
16 loss = -tf.reduce_sum(y * tf.log(probs))
17
18 learning_rate = 1e-2
19 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
20
21 xx = np.random.randn(N, D).astype(np.float32)
22 yy = np.zeros((N, C)).astype(np.float32)
23 yy[np.arange(N), np.random.randint(C, size=N)] = 1
24
25 with tf.Session() as sess:
26     sess.run(tf.initialize_all_variables())
27
28     for t in xrange(100):
29         _, loss_value = sess.run([train_step, loss],
30                               feed_dict={x: xx, y: yy})
31         print loss_value
```

# TensorFlow: Two-Layer Net

Create placeholders for data and labels: These will be fed to the graph

```
 1 import tensorflow as tf
 2 import numpy as np
 3
 4 N, D, H, C = 64, 1000, 100, 10
 5
 6 x = tf.placeholder(tf.float32, shape=[None, D])
 7 y = tf.placeholder(tf.float32, shape=[None, C])
 8
 9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32))
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32))
11
12 a = tf.matmul(x, w1)
13 a_relu = tf.nn.relu(a)
14 scores = tf.matmul(a_relu, w2)
15 probs = tf.nn.softmax(scores)
16 loss = -tf.reduce_sum(y * tf.log(probs))
17
18 learning_rate = 1e-2
19 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
20
21 xx = np.random.randn(N, D).astype(np.float32)
22 yy = np.zeros((N, C)).astype(np.float32)
23 yy[np.arange(N), np.random.randint(C, size=N)] = 1
24
25 with tf.Session() as sess:
26     sess.run(tf.initialize_all_variables())
27
28     for t in xrange(100):
29         _, loss_value = sess.run([train_step, loss],
30                               feed_dict={x: xx, y: yy})
31         print loss_value
```

# TensorFlow: Two-Layer Net

Create Variables to hold weights; similar to Theano shared variables

Initialize variables with numpy arrays

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D])
7 y = tf.placeholder(tf.float32, shape=[None, C])
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32))
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32))
11
12 a = tf.matmul(x, w1)
13 a_relu = tf.nn.relu(a)
14 scores = tf.matmul(a_relu, w2)
15 probs = tf.nn.softmax(scores)
16 loss = -tf.reduce_sum(y * tf.log(probs))
17
18 learning_rate = 1e-2
19 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
20
21 xx = np.random.randn(N, D).astype(np.float32)
22 yy = np.zeros((N, C)).astype(np.float32)
23 yy[np.arange(N), np.random.randint(C, size=N)] = 1
24
25 with tf.Session() as sess:
26     sess.run(tf.initialize_all_variables())
27
28     for t in xrange(100):
29         loss_value = sess.run([train_step, loss],
30                             feed_dict={x: xx, y: yy})
31         print loss_value
```

# TensorFlow: Two-Layer Net

**Forward:** Compute scores,  
probs, loss (symbolically)

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D])
7 y = tf.placeholder(tf.float32, shape=[None, C])
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32))
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32))
11
12 a = tf.matmul(x, w1)
13 a_relu = tf.nn.relu(a)
14 scores = tf.matmul(a_relu, w2)
15 probs = tf.nn.softmax(scores)
16 loss = -tf.reduce_sum(y * tf.log(probs))
17
18 learning_rate = 1e-2
19 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
20
21 xx = np.random.randn(N, D).astype(np.float32)
22 yy = np.zeros((N, C)).astype(np.float32)
23 yy[np.arange(N), np.random.randint(C, size=N)] = 1
24
25 with tf.Session() as sess:
26     sess.run(tf.initialize_all_variables())
27
28     for t in xrange(100):
29         _, loss_value = sess.run([train_step, loss],
30                                 feed_dict={x: xx, y: yy})
31         print loss_value
```

# TensorFlow: Two-Layer Net

Running train\_step will  
use SGD to minimize loss

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D])
7 y = tf.placeholder(tf.float32, shape=[None, C])
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32))
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32))
11
12 a = tf.matmul(x, w1)
13 a_relu = tf.nn.relu(a)
14 scores = tf.matmul(a_relu, w2)
15 probs = tf.nn.softmax(scores)
16 loss = -tf.reduce_sum(y * tf.log(probs))
17
18 learning_rate = 1e-2
19 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
20
21 xx = np.random.randn(N, D).astype(np.float32)
22 yy = np.zeros((N, C)).astype(np.float32)
23 yy[np.arange(N), np.random.randint(C, size=N)] = 1
24
25 with tf.Session() as sess:
26     sess.run(tf.initialize_all_variables())
27
28     for t in xrange(100):
29         _, loss_value = sess.run([train_step, loss],
30                                feed_dict={x: xx, y: yy})
31         print loss_value
```

# TensorFlow: Two-Layer Net

Create an artificial dataset;  
y is one-hot like Keras

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D])
7 y = tf.placeholder(tf.float32, shape=[None, C])
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32))
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32))
11
12 a = tf.matmul(x, w1)
13 a_relu = tf.nn.relu(a)
14 scores = tf.matmul(a_relu, w2)
15 probs = tf.nn.softmax(scores)
16 loss = -tf.reduce_sum(y * tf.log(probs))
17
18 learning_rate = 1e-2
19 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
20
21 xx = np.random.randn(N, D).astype(np.float32)
22 yy = np.zeros((N, C)).astype(np.float32)
23 yy[np.arange(N), np.random.randint(C, size=N)] = 1
24
25 with tf.Session() as sess:
26     sess.run(tf.initialize_all_variables())
27
28     for t in xrange(100):
29         _, loss_value = sess.run([train_step, loss],
30                                feed_dict={x: xx, y: yy})
31         print loss_value
```



# TensorFlow: Two-Layer Net

Actually train the model

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D])
7 y = tf.placeholder(tf.float32, shape=[None, C])
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32))
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32))
11
12 a = tf.matmul(x, w1)
13 a_relu = tf.nn.relu(a)
14 scores = tf.matmul(a_relu, w2)
15 probs = tf.nn.softmax(scores)
16 loss = -tf.reduce_sum(y * tf.log(probs))
17
18 learning_rate = 1e-2
19 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
20
21 xx = np.random.randn(N, D).astype(np.float32)
22 yy = np.zeros((N, C)).astype(np.float32)
23 yy[np.arange(N), np.random.randint(C, size=N)] = 1
24
25 with tf.Session() as sess:
26     sess.run(tf.initialize_all_variables())
27
28     for t in xrange(100):
29         loss_value = sess.run([train_step, loss],
30                             feed_dict={x: xx, y: yy})
31         print loss_value
```

# TensorFlow: Tensorboard

Tensorboard makes it easy to visualize what's happening inside your models

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D])
7 y = tf.placeholder(tf.float32, shape=[None, C])
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32))
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32))
11
12 a = tf.matmul(x, w1)
13 a_relu = tf.nn.relu(a)
14 scores = tf.matmul(a_relu, w2)
15 probs = tf.nn.softmax(scores)
16 loss = -tf.reduce_sum(y * tf.log(probs))
17
18 loss_summary = tf.scalar_summary('loss', loss)
19 w1_hist = tf.histogram_summary('w1', w1)
20 w2_hist = tf.histogram_summary('w2', w2)
21
22 learning_rate = 1e-2
23 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
24
25 xx = np.random.randn(N, D).astype(np.float32)
26 yy = np.zeros((N, C)).astype(np.float32)
27 yy[np.arange(N), np.random.randint(C, size=N)] = 1
28
29 with tf.Session() as sess:
30     merged = tf.merge_all_summaries()
31     writer = tf.train.SummaryWriter('/tmp/fc_logs', sess.graph_def)
32     sess.run(tf.initialize_all_variables())
33
34     for t in xrange(100):
35         summary_str, _, loss_value = sess.run(
36             [merged, train_step, loss],
37             feed_dict={x: xx, y: yy})
38         writer.add_summary(summary_str, t)
39         print loss_value
```

# TensorFlow: Tensorboard

Tensorboard makes it easy to visualize what's happening inside your models

Same as before, but now we create summaries for loss and weights

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D])
7 y = tf.placeholder(tf.float32, shape=[None, C])
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32))
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32))
11
12 a = tf.matmul(x, w1)
13 a_relu = tf.nn.relu(a)
14 scores = tf.matmul(a_relu, w2)
15 probs = tf.nn.softmax(scores)
16 loss = -tf.reduce_sum(y * tf.log(probs))
17
18 loss_summary = tf.scalar_summary('loss', loss)
19 w1_hist = tf.histogram_summary('w1', w1)
20 w2_hist = tf.histogram_summary('w2', w2)
21
22 learning_rate = 1e-2
23 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
24
25 xx = np.random.randn(N, D).astype(np.float32)
26 yy = np.zeros((N, C)).astype(np.float32)
27 yy[np.arange(N), np.random.randint(C, size=N)] = 1
28
29 with tf.Session() as sess:
30     merged = tf.merge_all_summaries()
31     writer = tf.train.SummaryWriter('/tmp/fc_logs', sess.graph_def)
32     sess.run(tf.initialize_all_variables())
33
34     for t in xrange(100):
35         summary_str, _, loss_value = sess.run(
36             [merged, train_step, loss],
37             feed_dict={x: xx, y: yy})
38         writer.add_summary(summary_str, t)
39         print loss_value
```

# TensorFlow: Tensorboard

Tensorboard makes it easy to visualize what's happening inside your models

Create a special “merged” variable and a SummaryWriter object

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D])
7 y = tf.placeholder(tf.float32, shape=[None, C])
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32))
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32))
11
12 a = tf.matmul(x, w1)
13 a_relu = tf.nn.relu(a)
14 scores = tf.matmul(a_relu, w2)
15 probs = tf.nn.softmax(scores)
16 loss = -tf.reduce_sum(y * tf.log(probs))
17
18 loss_summary = tf.scalar_summary('loss', loss)
19 w1_hist = tf.histogram_summary('w1', w1)
20 w2_hist = tf.histogram_summary('w2', w2)
21
22 learning_rate = 1e-2
23 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
24
25 xx = np.random.randn(N, D).astype(np.float32)
26 yy = np.zeros((N, C)).astype(np.float32)
27 yy[np.arange(N), np.random.randint(C, size=N)] = 1
28
29 with tf.Session() as sess:
30     merged = tf.merge_all_summaries()
31     writer = tf.train.SummaryWriter('/tmp/fc_logs', sess.graph_def)
32     sess.run(tf.initialize_all_variables())
33
34     for t in xrange(100):
35         summary_str, _, loss_value = sess.run(
36             (merged, train_step, loss),
37             feed_dict={x: xx, y: yy})
38         writer.add_summary(summary_str, t)
39         print loss_value
```

# TensorFlow: Tensorboard

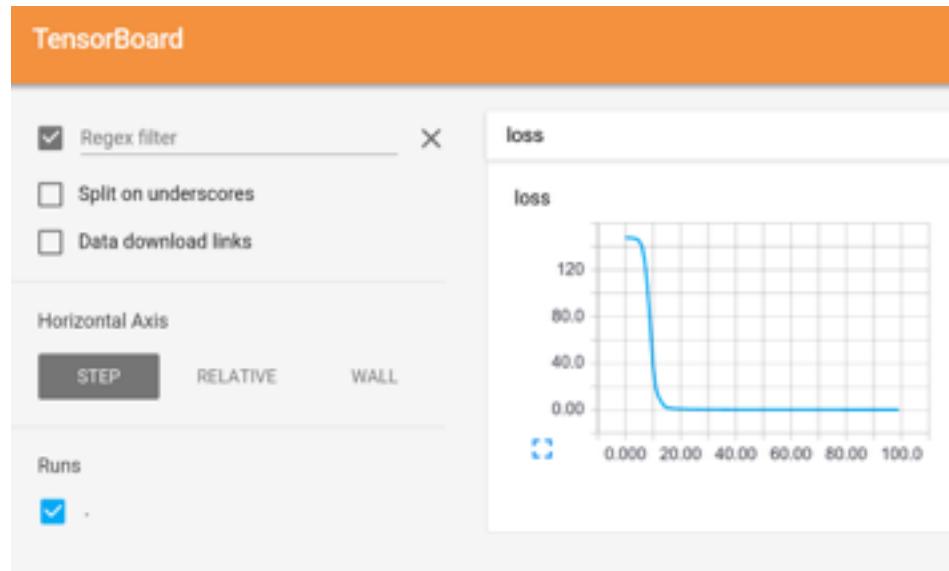
Tensorboard makes it easy to visualize what's happening inside your models

In the training loop, also run merged and pass its value to the writer

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D])
7 y = tf.placeholder(tf.float32, shape=[None, C])
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32))
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32))
11
12 a = tf.matmul(x, w1)
13 a_relu = tf.nn.relu(a)
14 scores = tf.matmul(a_relu, w2)
15 probs = tf.nn.softmax(scores)
16 loss = -tf.reduce_sum(y * tf.log(probs))
17
18 loss_summary = tf.scalar_summary('loss', loss)
19 w1_hist = tf.histogram_summary('w1', w1)
20 w2_hist = tf.histogram_summary('w2', w2)
21
22 learning_rate = 1e-2
23 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
24
25 xx = np.random.randn(N, D).astype(np.float32)
26 yy = np.zeros((N, C)).astype(np.float32)
27 yy[np.arange(N), np.random.randint(C, size=N)] = 1
28
29 with tf.Session() as sess:
30     merged = tf.merge_all_summaries()
31     writer = tf.train.SummaryWriter('/tmp/fc_logs', sess.graph_def)
32     sess.run(tf.initialize_all_variables())
33
34     for t in xrange(100):
35         summary_str, _, loss_value = sess.run(
36             (merged, train_step, loss),
37             feed_dict={x: xx, y: yy})
38         writer.add_summary(summary_str, t)
39         print loss_value
```



# TensorFlow: Tensorboard



Start Tensorboard server, and we get graphs!

# TensorFlow: TensorBoard

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D], name='x')
7 y = tf.placeholder(tf.float32, shape=[None, C], name='y')
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32), name='w1')
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32), name='w2')
11
12 with tf.name_scope('scores') as scope:
13     a = tf.matmul(x, w1)
14     a_relu = tf.nn.relu(a)
15     scores = tf.matmul(a_relu, w2)
16 with tf.name_scope('loss') as scope:
17     probs = tf.nn.softmax(scores)
18     loss = -tf.reduce_sum(y * tf.log(probs))
19
20 loss_summary = tf.scalar_summary('loss', loss)
21 w1_hist = tf.histogram_summary('w1', w1)
22 w2_hist = tf.histogram_summary('w2', w2)
23
24 learning_rate = 1e-2
25 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
26
27 xx = np.random.randn(N, D).astype(np.float32)
28 yy = np.zeros((N, C)).astype(np.float32)
29 yy[np.arange(N), np.random.randint(C, size=N)] = 1
30
```

# TensorFlow: TensorBoard

Add names to placeholders  
and variables

```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D], name='x')
7 y = tf.placeholder(tf.float32, shape=[None, C], name='y')
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32), name='w1')
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32), name='w2')
11
12 with tf.name_scope("scores") as scope:
13     a = tf.matmul(x, w1)
14     a_relu = tf.nn.relu(a)
15     scores = tf.matmul(a_relu, w2)
16 with tf.name_scope("loss") as scope:
17     probs = tf.nn.softmax(scores)
18     loss = -tf.reduce_sum(y * tf.log(probs))
19
20 loss_summary = tf.scalar_summary('loss', loss)
21 w1_hist = tf.histogram_summary('w1', w1)
22 w2_hist = tf.histogram_summary('w2', w2)
23
24 learning_rate = 1e-2
25 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
26
27 xx = np.random.randn(N, D).astype(np.float32)
28 yy = np.zeros((N, C)).astype(np.float32)
29 yy[np.arange(N), np.random.randint(C, size=N)] = 1
```

# TensorFlow: TensorBoard

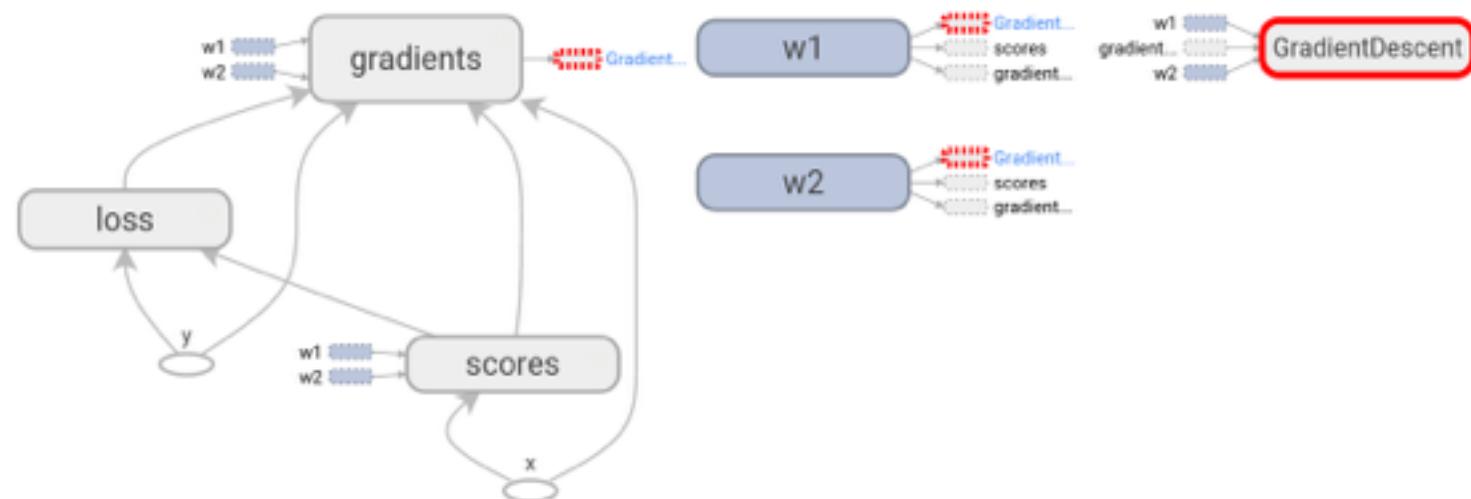
Add names to placeholders and variables

Break up the forward pass with name scoping



```
1 import tensorflow as tf
2 import numpy as np
3
4 N, D, H, C = 64, 1000, 100, 10
5
6 x = tf.placeholder(tf.float32, shape=[None, D], name='x')
7 y = tf.placeholder(tf.float32, shape=[None, C], name='y')
8
9 w1 = tf.Variable(1e-3 * np.random.randn(D, H).astype(np.float32), name='w1')
10 w2 = tf.Variable(1e-3 * np.random.randn(H, C).astype(np.float32), name='w2')
11
12 with tf.name_scope('scores') as scope:
13     a = tf.matmul(x, w1)
14     a_relu = tf.nn.relu(a)
15     scores = tf.matmul(a_relu, w2)
16 with tf.name_scope('loss') as scope:
17     probs = tf.nn.softmax(scores)
18     loss = -tf.reduce_sum(y * tf.log(probs))
19
20 loss_summary = tf.scalar_summary('loss', loss)
21 w1_hist = tf.histogram_summary('w1', w1)
22 w2_hist = tf.histogram_summary('w2', w2)
23
24 learning_rate = 1e-2
25 train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)
26
27 xx = np.random.randn(N, D).astype(np.float32)
28 yy = np.zeros((N, C)).astype(np.float32)
29 yy[np.arange(N), np.random.randint(C, size=N)] = 1
```

# TensorFlow: TensorBoard

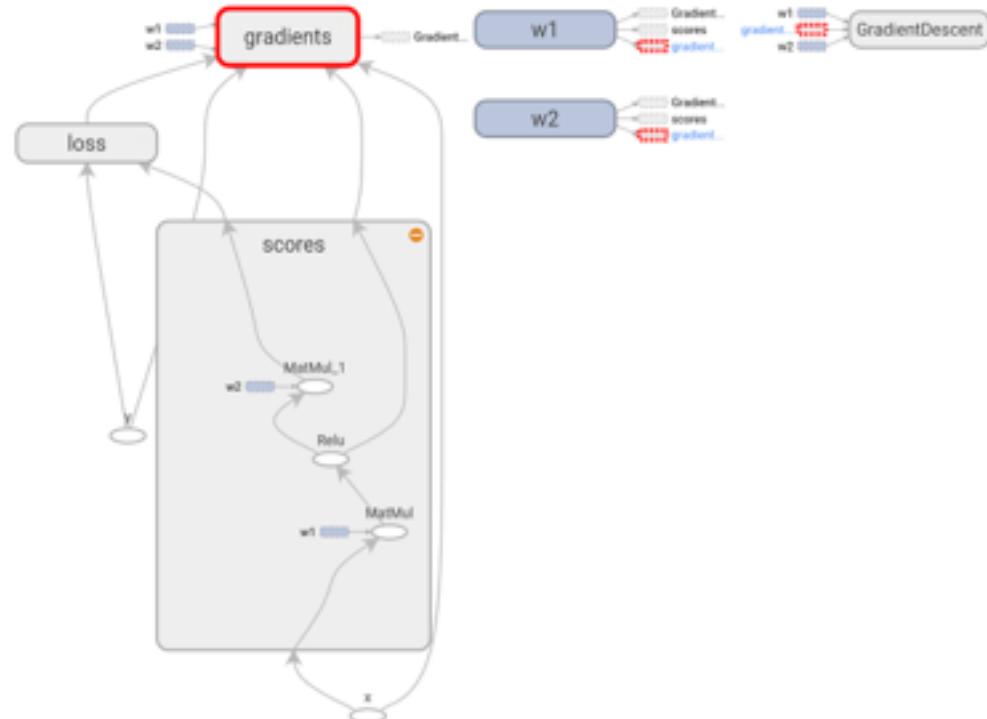


Tensorboard shows the graph!

# TensorFlow: TensorBoard

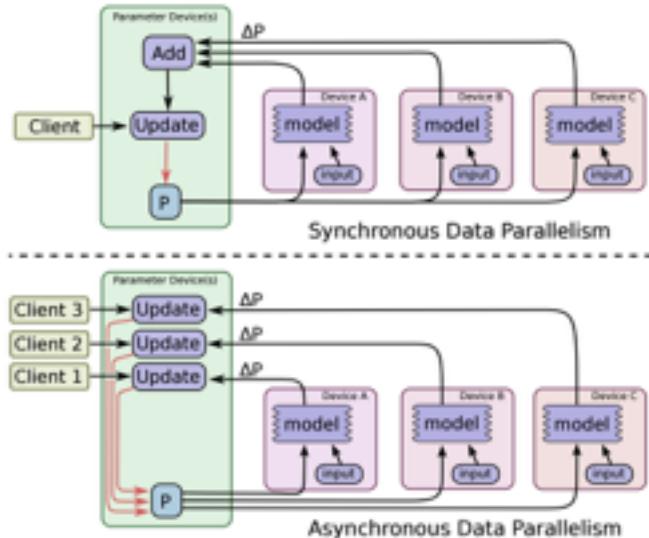
Tensorboard shows the graph!

Name scopes expand to show individual operations



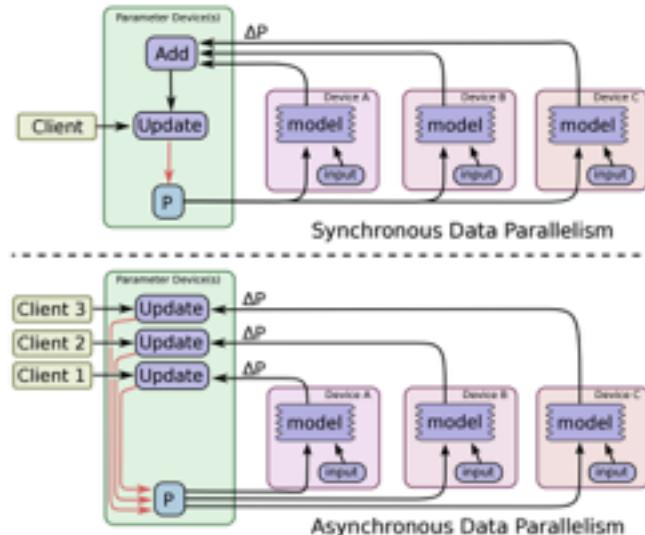
# TensorFlow: Multi-GPU

**Data parallelism:**  
synchronous or asynchronous

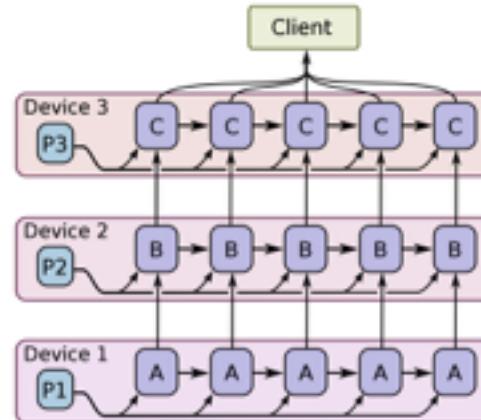


# TensorFlow: Multi-GPU

**Data parallelism:**  
synchronous or asynchronous



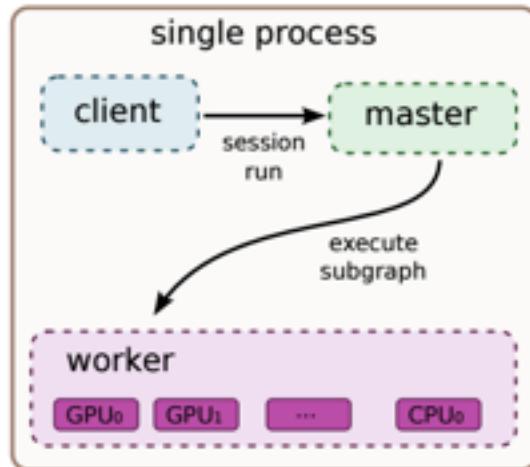
**Model parallelism:**  
Split model across GPUs



# TensorFlow: Distributed

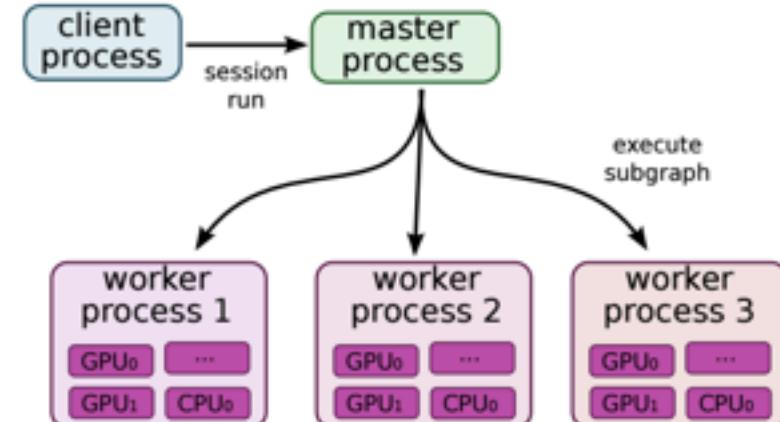
## Single machine:

Like other frameworks



## Many machines:

Not open source (yet) =(



# TensorFlow: Pretrained Models

- You can get a pretrained models here:
  - <https://github.com/tensorflow/models>
- Has inception, resnet, some different autoencoders

# TensorFlow: Pros / Cons

- (+) Python + numpy
- (+) Computational graph abstraction, like Theano; great for RNNs
- (+) Much faster compile times than Theano
- (+) Slightly more convenient than raw Theano?
- (+) TensorBoard for visualization
- (+) Data AND model parallelism; best of all frameworks
- (+/-) Distributed models, but not open-source yet
- (-) Slower than other frameworks right now
- (-) Much “fatter” than Torch; more magic
- (-) Not many pretrained models

# Torch

<http://torch.ch/docs/getting-started.html>

# Torch Overview

- From NYU + IDIAP
- Written in C and Lua
- Used a lot at Facebook, DeepMind

# Torch: Pretrained Models

- **loadcaffe**: Load pretrained Caffe models: AlexNet, VGG, some others  
<https://github.com/szagoruyko/loadcaffe>
- **GoogLeNet v1**: <https://github.com/soumith/inception.torch>
- **GoogLeNet v3**: <https://github.com/Moodstocks/inception-v3.torch>
- **ResNet**: <https://github.com/facebook/fb.resnet.torch>

# Torch: Pros / Cons

- (-) Lua
- (-) Less plug-and-play than Caffe
  - You usually write your own training code
- (+) Lots of modular pieces that are easy to combine
- (+) Easy to write your own layer types and run on GPU
- (+) Most of the library code is in Lua, easy to read
- (+) Lots of pretrained models!
- (-) Not great for RNNs

# Theano

<http://deeplearning.net/software/theano/>

# Theano Overview

- From Yoshua Bengio's group at University of Montreal
- Embracing computation graphs, symbolic computation
- High-level wrappers: Keras, Lasagne
- Has Conditional flow (ifelse, switch)

# Theano: Pretrained Models

Best choice

- **Lasagne Model Zoo** has pretrained common architectures:  
<https://github.com/Lasagne/Recipes/tree/master/modelzoo>
- **AlexNet with weights**: [https://github.com/uoguelph-mlrg/theano\\_alexnet](https://github.com/uoguelph-mlrg/theano_alexnet)
- **sklearn-theano**: Run OverFeat and GoogLeNet forward, but no fine-tuning? <http://sklearn-theano.github.io>
- **caffe-theano-conversion**: CS 231n project from last year: load models and weights from caffe! Not sure if full-featured <https://github.com/kitofans/caffe-theano-conversion>

# Theano: Pros / Cons

- (+) Python + numpy
- (+) Computational graph is nice abstraction
- (+) RNNs fit nicely in computational graph
- (-) Raw Theano is somewhat low-level
- (+) High level wrappers (Keras, Lasagne) ease the pain
- (-) Error messages can be unhelpful
- (-) Large models can have long compile times
- (-) Much “fatter” than Torch; more magic
- (-) Patchy support for pretrained models

# Overview

	Caffe	Torch	Theano	TensorFlow
<b>Language</b>	C++, Python	Lua	Python	Python
<b>Pretrained</b>	Yes ++	Yes ++	Yes (Lasagne)	Inception
<b>Multi-GPU: Data parallel</b>	Yes	Yes cunn.DataParallelTable	Yes platoon	Yes
<b>Multi-GPU: Model parallel</b>	No	Yes fbcunn.ModelParallel	Experimental	Yes (best)
<b>Readable source code</b>	Yes (C++)	Yes (Lua)	No	No
<b>Good at RNN</b>	No	Mediocre	Yes	Yes (best)