

JOSEPH

ROSS

SANTOSH

ALI

REDMON

GIRSHICK

DIVVALA

FARHADI

Dog



“YOU ONLY LOOK ONCE”
**REAL-TIME
DETECTION**

Presented by:

Jie Li
Takuto Sato

Original Slides from:
<https://pjreddie.com/darknet/yolo/>


Classification + Localization



CAT

Object Detection

v.s.



CAT, DOG, DUCK

Localization as Regression

Input: image



Neural Net
→

Output:
Box coordinates
(4 numbers)

→ **Loss:**
L2 distance

Correct output:
box coordinates
(4 numbers)

Only one object,
simpler than detection


* Original slides borrowed from Andrej Karpathy
and Li Fei-Fei, Stanford cs231n

Sliding Window: Overfeat

(DPM as well!)

In practice use many sliding window locations and multiple scales

Window positions + score maps



Box regression outputs



Final Predictions



Sermanet et al, "Integrated Recognition, Localization and Detection using Convolutional Networks", ICLR 2014

* Original slides borrowed from Andrej Karpathy and Li Fei-Fei, Stanford cs231n

Detection as Regression?



DOG, (x, y, w, h)
CAT, (x, y, w, h)
CAT, (x, y, w, h)
DUCK (x, y, w, h)

= 16 numbers

Detection as Regression?



DOG, (x, y, w, h)
CAT, (x, y, w, h)
= 8 numbers

Detection as Regression?



CAT, (x, y, w, h)

CAT, (x, y, w, h)

....

CAT (x, y, w, h)

= many numbers

Need variable sized outputs

Detection as Classification



CAT? NO

DOG? NO


Detection as Classification



CAT? YES!

DOG? NO

Detection as Classification



CAT? NO

DOG? NO

Detection as Classification

Problem: Need to test many positions and scales

Solution: If your classifier is fast enough, just do it


Detection as Classification

Problem: Need to test many positions and scales,
and use a computationally demanding classifier (CNN)

Solution: Only look at a tiny subset of possible positions

Region Proposals


- Find “blobby” image regions that are likely to contain objects
- “Class-agnostic” object detector
- Look for “blob-like” regions



Region Proposals: Selective Search

Bottom-up segmentation, merging regions at multiple scales

Convert
regions
to boxes




Uijlings et al, "Selective Search for Object Recognition", IJCV 2013

Region Proposals: Many other choices

Method	Approach	Outputs Segments	Outputs Score	Control #proposals	Time (sec.)	Repeatability	Recall Results	Detection Results
Bing [18]	Window scoring		✓	✓	0.2	***	*	-
CPMC [19]	Grouping	✓	✓	✓	250	-	**	*
EdgeBoxes [20]	Window scoring		✓	✓	0.3	**	***	***
Endres [21]	Grouping	✓	✓	✓	100	-	***	**
Geodesic [22]	Grouping	✓		✓	1	*	***	**
MCG [23]	Grouping	✓	✓	✓	30	*	***	***
Objectness [24]	Window scoring		✓	✓	3	-	*	-
Rahtu [25]	Window scoring		✓	✓	3	-	-	*
RandomizedPrim's [26]	Grouping	✓		✓	1	*	*	**
Rantalankila [27]	Grouping	✓		✓	10	**	-	**
Rigor [28]	Grouping	✓		✓	10	*	**	**
SelectiveSearch [29]	Grouping	✓	✓	✓	10	**	***	***
Gaussian				✓	0	-	-	*
SlidingWindow				✓	0	***	-	-
Superpixels		✓			1	*	-	-
Uniform				✓	0	-	-	-

Hosang et al, "What makes for effective detection proposals?", PAMI 2015


Putting it together: R-CNN



Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014

Slide credit: Ross Girshick


Fast R-CNN (test time)



R-CNN Problem #1:
Slow at test-time due to
independent forward
passes of the CNN

Solution:
Share computation
of convolutional
layers between
proposals for an
image

Faster R-CNN:



Insert a **Region Proposal**

Network (RPN) after the last convolutional layer

RPN trained to produce region proposals directly; no need for external region proposals!

After RPN, use RoI Pooling and an upstream classifier and bbox regressor just like Fast R-CNN

Ren et al, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks", NIPS 2015

Slide credit: Ross Girshick

Summary:

How to frame detection problem?

Regression?

- OK for one or a fixed number of objects (localization with sliding windows)
- Sliding windows is slow (especially when the number of objects too large)
- Number of objects unknown?

Classification?

- Need region-based method to propose bounding boxes (~2000 boxes, R-CNN)
- Still kind of slow (especially considering real-time detection)

You Only Look Once: Unified, Real-Time Object Detection

Joseph Redmon*, Santosh Divvala*†, Ross Girshick¶, Ali Farhadi*†

University of Washington*, Allen Institute for AI†, Facebook AI Research¶

<http://pjreddie.com/yolo/>

Abstract

We present YOLO, a new approach to object detection. Prior work on object detection repurposes classifiers to perform detection. Instead, we frame object detection as a regression problem to spatially separated bounding boxes and associated class probabilities. A single neural network predicts bounding boxes and class probabilities directly from full images in one evaluation. Since the whole detection pipeline is a single network, it can be optimized end-to-end directly on detection performance.

Our unified architecture is extremely fast. Our base YOLO model processes images in real-time at 45 frames per second. A smaller version of the network, Fast YOLO, processes an astounding 155 frames per second while still achieving double the mAP of other real-time detectors. Compared to state-of-the-art detection systems, YOLO makes more localization errors but is less likely to predict false positives on background. Finally, YOLO learns very general representations of objects. It outperforms other de-




Figure 1: The YOLO Detection System. Processing images with YOLO is simple and straightforward. Our system (1) resizes the input image to 448×448 , (2) runs a single convolutional network on the image, and (3) thresholds the resulting detections by the model’s confidence.


methods to first generate potential bounding boxes in an image and then run a classifier on these proposed boxes. After classification, post-processing is used to refine the bounding boxes, eliminate duplicate detections, and rescore the boxes based on other objects in the scene [13]. These complex pipelines are slow and hard to optimize because each individual component must be trained separately.

We reframe object detection as a single regression prob-



Original Slides from:
<https://pjreddie.com/darknet/yolo/>

We split the image into a grid




$$S = 7$$

Total # of Cells:
 $S \times S = 49$

Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Each cell predicts boxes and confidences: $P(\text{Object})$



$$S = 7$$

Total # of Cells:
 $S \times S = 49$

Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Each cell predicts boxes and confidences: $P(\text{Object})$




$$S = 7$$

$$\begin{aligned} \text{Total # of Cells:} \\ S \times S = 49 \end{aligned}$$

$$B = 2$$

Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Each cell predicts boxes and confidences: $P(\text{Object})$




$$S = 7$$

$$\begin{aligned} \text{Total \# of Cells:} \\ S \times S = 49 \end{aligned}$$

$$B = 2$$

Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Each cell predicts boxes and confidences: $P(\text{Object})$




$$S = 7$$

$$\begin{aligned} \text{Total \# of Cells:} \\ S \times S = 49 \end{aligned}$$

$$B = 2$$

Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Each cell predicts boxes and confidences: $P(\text{Object})$




$$S = 7$$

$$\begin{aligned} \text{Total # of Cells:} \\ S \times S = 49 \end{aligned}$$

$$B = 2$$

Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Each cell predicts boxes and confidences: $P(\text{Object})$



$$S = 7$$


$$\begin{aligned} \text{Total # of Cells:} \\ S \times S = 49 \end{aligned}$$

$$B = 2$$

$$\begin{aligned} \text{Total # of boxes:} \\ S \times S \times B = 98 \end{aligned}$$

Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Each cell also predicts a class probability.



Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Each cell also predicts a class probability.

Bicycle

Car

Dog

Dining
Table



Original Slides from:
<https://pjreddie.com/darknet/yolo/>


Conditioned on object: $P(\text{Car} \mid \text{Object})$

Bicycle


Car

Dog

Dining
Table




Then we combine the box and class predictions.



Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Finally we do NMS and threshold detections




Original Slides from:
<https://pjreddie.com/darknet/yolo/>

This parameterization fixes the output size

Each cell predicts:

- For each bounding box:
 - 4 coordinates (x, y, w, h)
 - 1 confidence value
- Some number of class probabilities




For Pascal VOC:

- 7x7 grid
- 2 bounding boxes / cell
- 20 classes

$$7 \times 7 \times (2 \times 5 + 20) = 7 \times 7 \times 30 \text{ tensor} = \mathbf{1470 \text{ outputs}}$$


Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Thus we can train one neural network to be a whole detection pipeline




Original Slides from:
<https://pjreddie.com/darknet/yolo/>

During training, match example to the right cell




Original Slides from:
<https://pjreddie.com/darknet/yolo/>

During training, match example to the right cell



Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Adjust that cell's class prediction



Dog = 1


Cat = 0

Bike = 0

...


Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Look at that cell's predicted boxes




Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Find the best one, make it “responsible for that prediction



Original Slides from:
<https://pjreddie.com/darknet/yolo/>

Find the best one, make it “responsible for that prediction



Original Slides from:
<https://pjreddie.com/darknet/yolo/>

loss function:

$$\begin{aligned}
& \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2 \right] \\
& + \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[\left(\sqrt{w_i} - \sqrt{\hat{w}_i} \right)^2 + \left(\sqrt{h_i} - \sqrt{\hat{h}_i} \right)^2 \right] \\
& + \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left(C_i - \hat{C}_i \right)^2 \\
& + \lambda_{\text{noobj}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{noobj}} \left(C_i - \hat{C}_i \right)^2 \\
& + \sum_{i=0}^{S^2} \mathbb{1}_i^{\text{obj}} \sum_{c \in \text{classes}} \left(p_i(c) - \hat{p}_i(c) \right)^2 \quad (3)
\end{aligned}$$

Strengths and Weaknesses (compared to the R-CNN family)

Strengths

- Fast
- Generalizes better (i.e. good at detecting objects in paintings)
- Less background false positives

Weaknesses

- Localization not as accurate
- Limited to B objects per grid ($B=2$ in paper/this presentation)




Figure 3: The Architecture. Our detection network has 24 convolutional layers followed by 2 fully connected layers. Alternating 1×1 convolutional layers reduce the features space from preceding layers. We pretrain the convolutional layers on the ImageNet classification task at half the resolution (224×224 input image) and then double the resolution for detection.